

Business Case: Scaler - Clustering

About the Dataset

- We are provided with the information for a segment of learners by Scaler, an online tech-versity.
- Working as a data scientist with the analytics vertical of Scaler, we got this dataset from the Scaler database.

Column Profiling:

1. 'Unnamed 0'- Index of the dataset
2. Email_hash- Anonymised Personal Identifiable Information (PII)
3. Company_hash- Current employer of the learner
4. orgyear- Employment start date
5. CTC- Current CTC
6. Job_position- Job profile in the company
7. CTC_updated_year: Year in which CTC got updated (Yearly increments, Promotions)

Problem Statement:

- We are tasked to cluster them on the basis of their job profile, company, and other features. We are focused on profiling the best companies and job positions to work.

Analysing basic metrics

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib
import matplotlib.pyplot as plt
from scipy import stats
import re
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: #Loading of dataset
df = pd.read_csv("../scaler/scaler_clustering.csv")
df.head()
```

Out[2]:

		Unnamed: 0	company_hash	email_hash	orgyear
0	0	atrgxnnt xzaxv	6de0a4417d18ab14334c3f43397fc13b30c35149d70c05...	201	
1	1	qtrxvzwt xzegwgbbr rxbxnta	b0aa1ac138b53cb6e039ba2c3d6604a250d02d5145c10...	201	
2	2	ojzwnvwnxw vx	4860c670bcd48fb96c02a4b0ae3608ae6fdd98176112e9...	201	
3	3	ngpgutaxv	effdede7a2e7c2af664c8a31d9346385016128d66bbc58...	201	
4	4	qxen sqghu	6ff54e709262f55cb999a1c1db8436cb2055d8f79ab520...	201	

In [3]: `df.shape #to observe shape of data`

Out[3]: (205843, 7)

- Dataset is of 205843 rows and 7 attributes.

In [4]: `df.info() #to observe the data type`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205843 entries, 0 to 205842
Data columns (total 7 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        205843 non-null   int64  
 1   company_hash     205799 non-null   object  
 2   email_hash       205843 non-null   object  
 3   orgyear          205757 non-null   float64 
 4   ctc              205843 non-null   int64  
 5   job_position     153281 non-null   object  
 6   ctc_updated_year 205843 non-null   float64 
dtypes: float64(2), int64(2), object(3)
memory usage: 11.0+ MB
```

Check for Duplicate Values

In [5]: `df.duplicated().sum()`

Out[5]: 0

- There are no duplicate instances in the data

Check for Missing Values

In [6]: `# Check Missing Values
df.isna().sum()`

```
Out[6]: Unnamed: 0          0
         company_hash      44
         email_hash        0
         orgyear          86
         ctc              0
         job_position     52562
         ctc_updated_year    0
         dtype: int64
```

- Data contains null values in 3 columns company_hash, orgyear, job_position.

```
In [7]: df.describe()
```

```
Out[7]:      Unnamed: 0       orgyear        ctc  ctc_updated_year
count  205843.000000  205757.000000  2.058430e+05  205843.000000
mean   103273.941786  2014.882750  2.271685e+06  2019.628231
std    59741.306484   63.571115  1.180091e+07  1.325104
min    0.000000     0.000000  2.000000e+00  2015.000000
25%   51518.500000  2013.000000  5.300000e+05  2019.000000
50%   103151.000000  2016.000000  9.500000e+05  2020.000000
75%   154992.500000  2018.000000  1.700000e+06  2021.000000
max   206922.000000  20165.000000  1.000150e+09  2021.000000
```

```
In [8]: df.describe(include="object")
```

```
Out[8]:      company_hash           email_hash  job_position
count          205799                  205843      153281
unique         37299                  153443      1017
top            nvnv
               wgzohrnvwj  bbace3cc586400bbc65765bc6a16b77d8913836cf98b7...
               otqcxwto
freq           8337                      10      43554
```

Data Preprocessing

```
In [9]: #Drop installment
df.drop(columns=['Unnamed: 0'], inplace=True)
```

```
In [10]: def preprocess_string(string):
    new_string= re.sub('[^A-Za-z ]+', ' ', string).lower().strip()
    return new_string

mystring='\tAirtel\\ \\&**() X Labs'
preprocess_string(mystring)
```

```
Out[10]: 'airtel x labs'
```

```
In [11]: df["company_hash"].nunique()
```

```
Out[11]: 37299
```

```
In [12]: df["company_hash"] = df["company_hash"].apply(lambda x: preprocess_string(str(x)))
df["company_hash"].nunique()
```

```
Out[12]: 37208
```

```
In [13]: df["job_position"].nunique()
```

```
Out[13]: 1017
```

- 1017 unique job positions are there in the dataset.

```
In [14]: df["job_position"] = df["job_position"].apply(lambda x: preprocess_string(str(x)))
df["job_position"].nunique()
```

```
Out[14]: 857
```

- 857 unique job positions are there in the dataset after preprocessing strings.

```
In [15]: df.isna().sum()
```

```
Out[15]: company_hash      0
email_hash        0
orgyear          86
ctc              0
job_position     0
ctc_updated_year 0
dtype: int64
```

```
In [16]: (df["company_hash"] == "").sum()
```

```
Out[16]: 89
```

```
In [17]: (df["company_hash"] == "nan").sum()
```

```
Out[17]: 44
```

```
In [18]: (df["job_position"] == "").sum()
```

```
Out[18]: 9
```

```
In [19]: (df["job_position"] == "nan").sum()
```

```
Out[19]: 52562
```

```
In [20]: # removing the records where company_hash or job_position records are not available
df[(df["company_hash"] == "") | (df["job_position"] == "")].sample(10)
```

Out[20]:

	company_hash	email_hash	orgyear
86378	77a5cecd2ed9bb764df8bf6da78a0ae2aef97fc87e913e...	2020.0	
187995	482bf44a26d01aa6e8313253f8956ea8cc2f2206be9307...	2019.0	
4151	f52c1fc6d5ee4c35ac8c7db5a8a5e0190b743d5487e288...	2020.0	
131273	df0652bad1b8a46bc63394e69e1b24d70ce092dc56243c...	2020.0	
156277	df0652bad1b8a46bc63394e69e1b24d70ce092dc56243c...	2020.0	
173726	6f5b3ae9ce591f3eafa4bfd89ce822ed80137601492621...	2018.0	
161047	yaew mvzp	0911dccca341fb4a54a729d0a5bf3adcc467c1ac0cf3322...	2001.0 1000000
4162		4793af2ce1550cdfd2e9ce20e234b321c1c180d46fe7f7...	2013.0
40476		7da9646f60f7e3e272deb8404fe92dee33909455fd83ea...	2018.0
22720		811c78a507b093aa1c736def6996fc3132fac1ad7d4840...	2019.0

In [21]:

```
len(df[(df["company_hash"] == "") | (df["job_position"] == "")])
```

Out[21]: 98

In [22]:

```
df = df[~((df["company_hash"] == "") | (df["job_position"] == ""))]
df.head()
```

Out[22]:

	company_hash	email_hash	orgyear	ctc
0	atrgxnnt xzaxv	6de0a4417d18ab14334c3f43397fc13b30c35149d70c05...	2016.0	1100000
1	qtrxvzwt xzegwgbbr rxbxnta	b0aa1ac138b53cb6e039ba2c3d6604a250d02d5145c10...	2018.0	449999
2	ojzwnvwnnxw vx	4860c670bcd48fb96c02a4b0ae3608ae6fd98176112e9...	2015.0	2000000
3	ngpgutaxv	effdede7a2e7c2af664c8a31d9346385016128d66bbc58...	2017.0	700000
4	qxen sqghu	6ff54e709262f55cb999a1c1db8436cb2055d8f79ab520...	2017.0	1400000

- Imputing Employee Start Year as per the median year as per each company.

In [23]:

```
df.groupby("company_hash")["orgyear"].transform("median")
```

```
Out[23]: 0      2014.0
1      2016.0
2      2015.0
3      2016.0
4      2017.0
...
205838 2018.0
205839 2017.0
205840 2016.0
205841 2020.0
205842 2015.0
Name: orgyear, Length: 205745, dtype: float64
```

```
In [24]: df["orgyear"].fillna(df['orgyear'].isnull().sum(), inplace=True)
```

```
In [25]: df["orgyear"].isna().sum()
```

```
Out[25]: 0
```

Non-Graphical Analysis

```
In [26]: # unique value company_hash column(Listed in %)
company_hash = df['company_hash'].value_counts(normalize=True).map(lambda calc:
company_hash.columns = ['company_hash', 'Count']
company_hash.head(10)
```

```
Out[26]:   company_hash  Count
0  nvnv wgzohrnvwj otqcxwto    4.05
1          xzegojo    2.62
2          vbvkgz    1.69
3  zgn vuurxwvmrt vwwghzn    1.66
4          wgszxkvzn    1.57
5          vwwwtzhqqt    1.39
6          fxuqg rxbxnta    1.29
7          gqvwrta    1.22
8          bxwqgogen    1.04
9          wvustbxzx    0.99
```

```
In [27]: # unique value orgyear column(Listed in %)
orgyear = df['orgyear'].value_counts(normalize=True).map(lambda calc: round(100*
orgyear.columns = ['orgyear', 'Count']
orgyear.head(10)
```

Out[27]:

	orgyear	Count
0	2018.0	12.27
1	2019.0	11.37
2	2017.0	11.29
3	2016.0	11.20
4	2015.0	10.01
5	2014.0	8.11
6	2020.0	6.52
7	2013.0	6.00
8	2012.0	5.10
9	2011.0	3.87

In [28]:

```
# unique value job_position column(listed in %)
job_position = df['job_position'].value_counts(normalize=True).map(lambda calc:
job_position.columns = ['job_position', 'Count']
job_position.head(10)
```

Out[28]:

	job_position	Count
0	nan	25.53
1	backend engineer	21.16
2	fullstack engineer	12.62
3	other	8.78
4	frontend engineer	5.06
5	engineering leadership	3.34
6	qa engineer	3.20
7	data scientist	2.61
8	android engineer	2.60
9	sdet	2.42

In [29]:

```
# unique value ctc_updated_year column(listed in %)
ctc_updated_year = df['ctc_updated_year'].value_counts(normalize=True).map(lambda calc:
ctc_updated_year.columns = ['ctc_updated_year', 'Count']
ctc_updated_year.head(10)
```

```
Out[29]:
```

	ctc_updated_year	Count
0	2019.0	33.37
1	2021.0	31.57
2	2020.0	24.02
3	2017.0	3.67
4	2018.0	3.28
5	2016.0	2.67
6	2015.0	1.42

```
In [30]:
```

```
# Number of unique values in all columns
unique_num = ['company_hash','email_hash','orgyear','ctc','job_position','ctc_updated_year']
for col in unique_num:
    print(f"No. of unique values in {col}: {df[col].nunique()}")
```

```
No. of unique values in company_hash: 37205
No. of unique values in email_hash: 153381
No. of unique values in orgyear: 78
No. of unique values in ctc: 3360
No. of unique values in job_position: 856
No. of unique values in ctc_updated_year: 7
```

```
In [31]:
```

```
df.describe().loc[['min', 'max']]
```

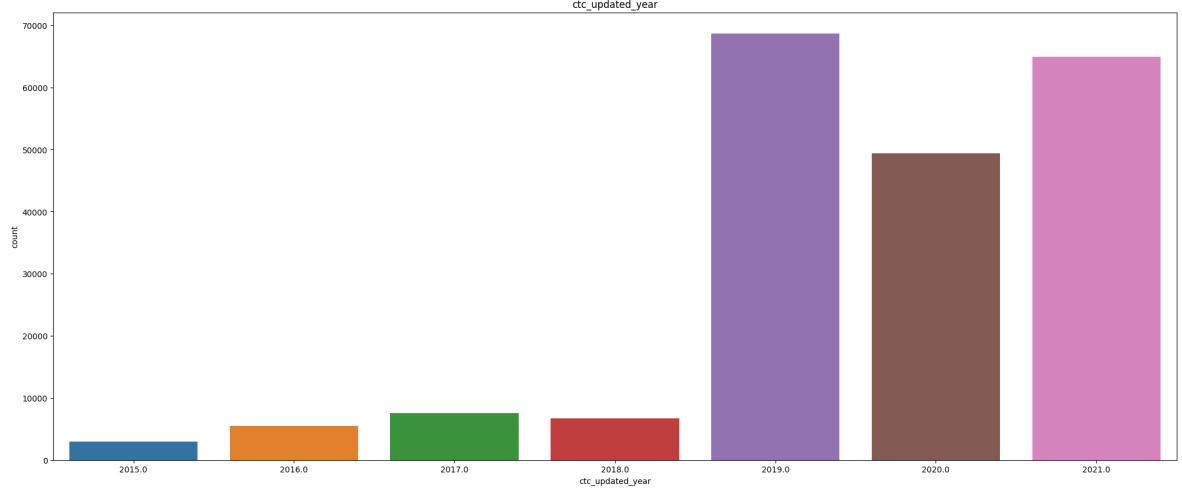
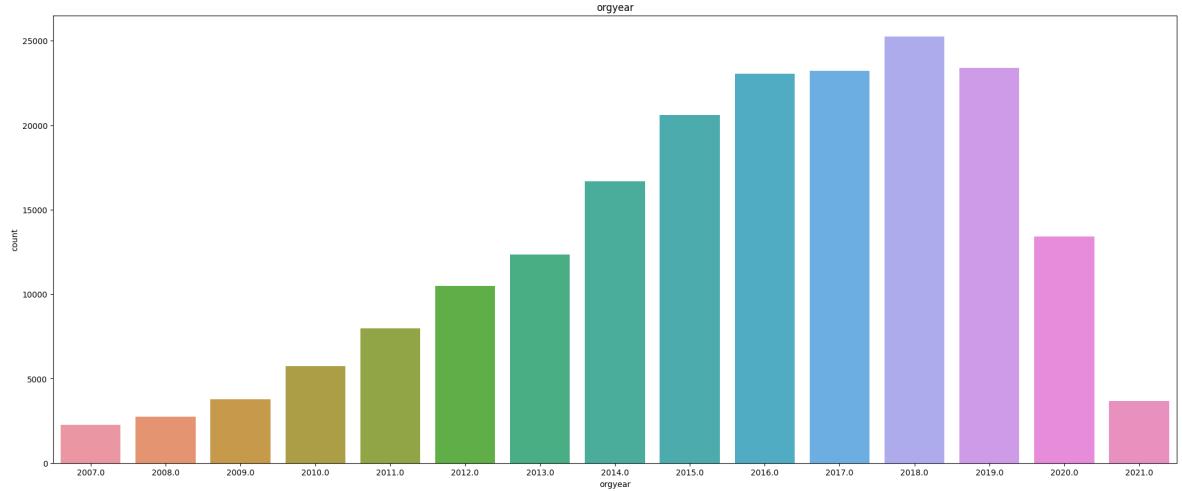
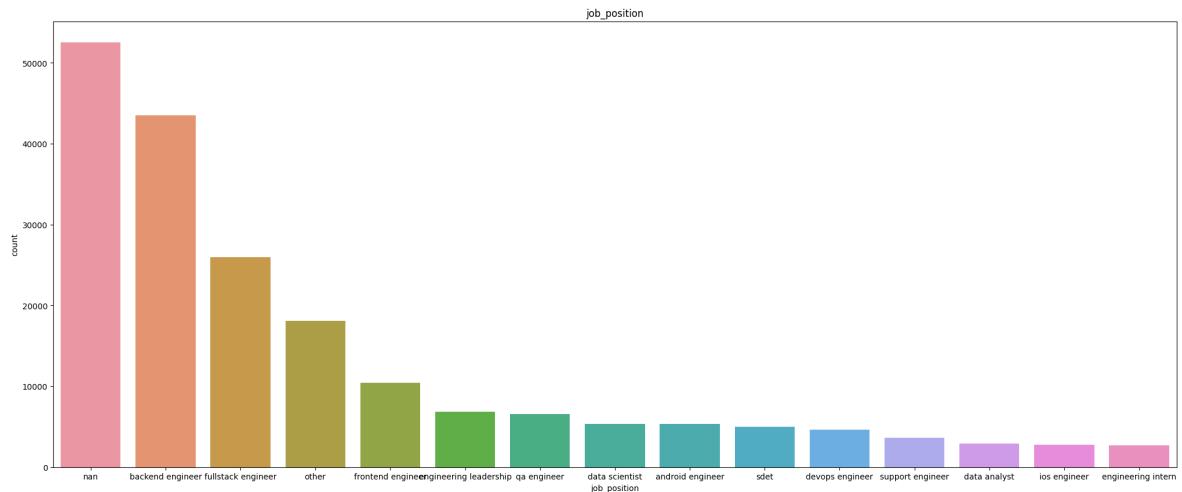
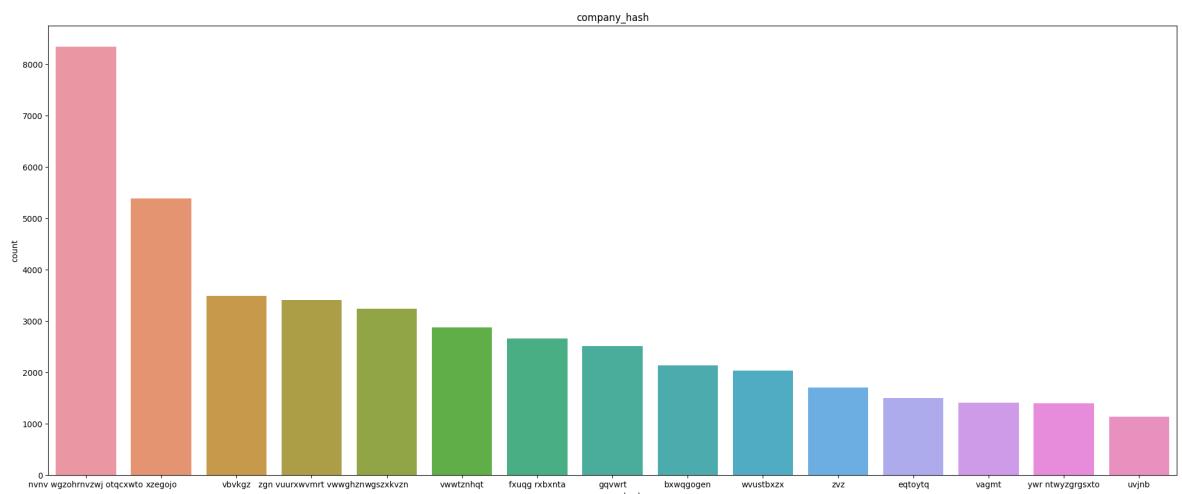
```
Out[31]:
```

	orgyear	ctc	ctc_updated_year
min	0.0	2.000000e+00	2015.0
max	20165.0	1.000150e+09	2021.0

Univariate Analysis

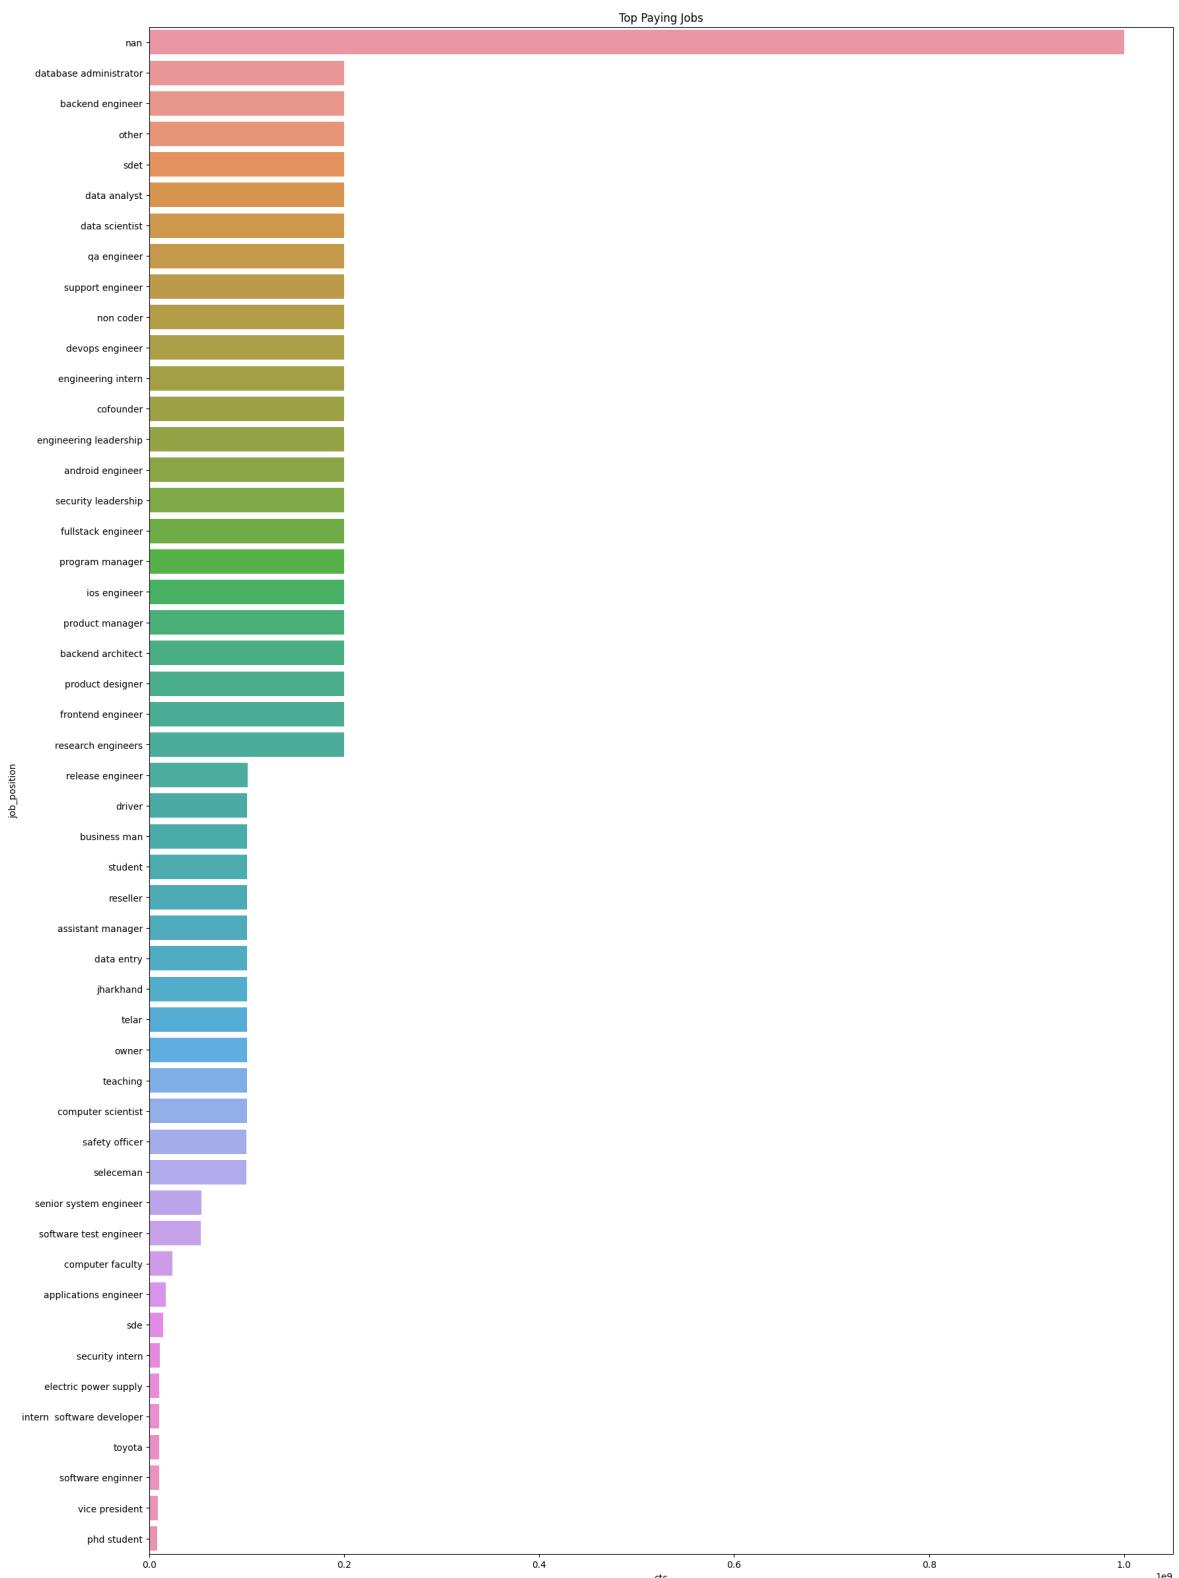
```
In [32]:
```

```
plot = ['company_hash','job_position','orgyear','ctc_updated_year']
for i in plot:
    tmp = df.copy()
    tmp['count'] = 1
    tmp = tmp.groupby(i).sum()['count'].reset_index().sort_values('count', ascending=False)
    plt.figure(figsize=(25,10))
    sns.barplot(data=tmp,y='count',x=i).set(title=i)
    plt.show()
```



Bivariate Analysis & Multivariate Analysis

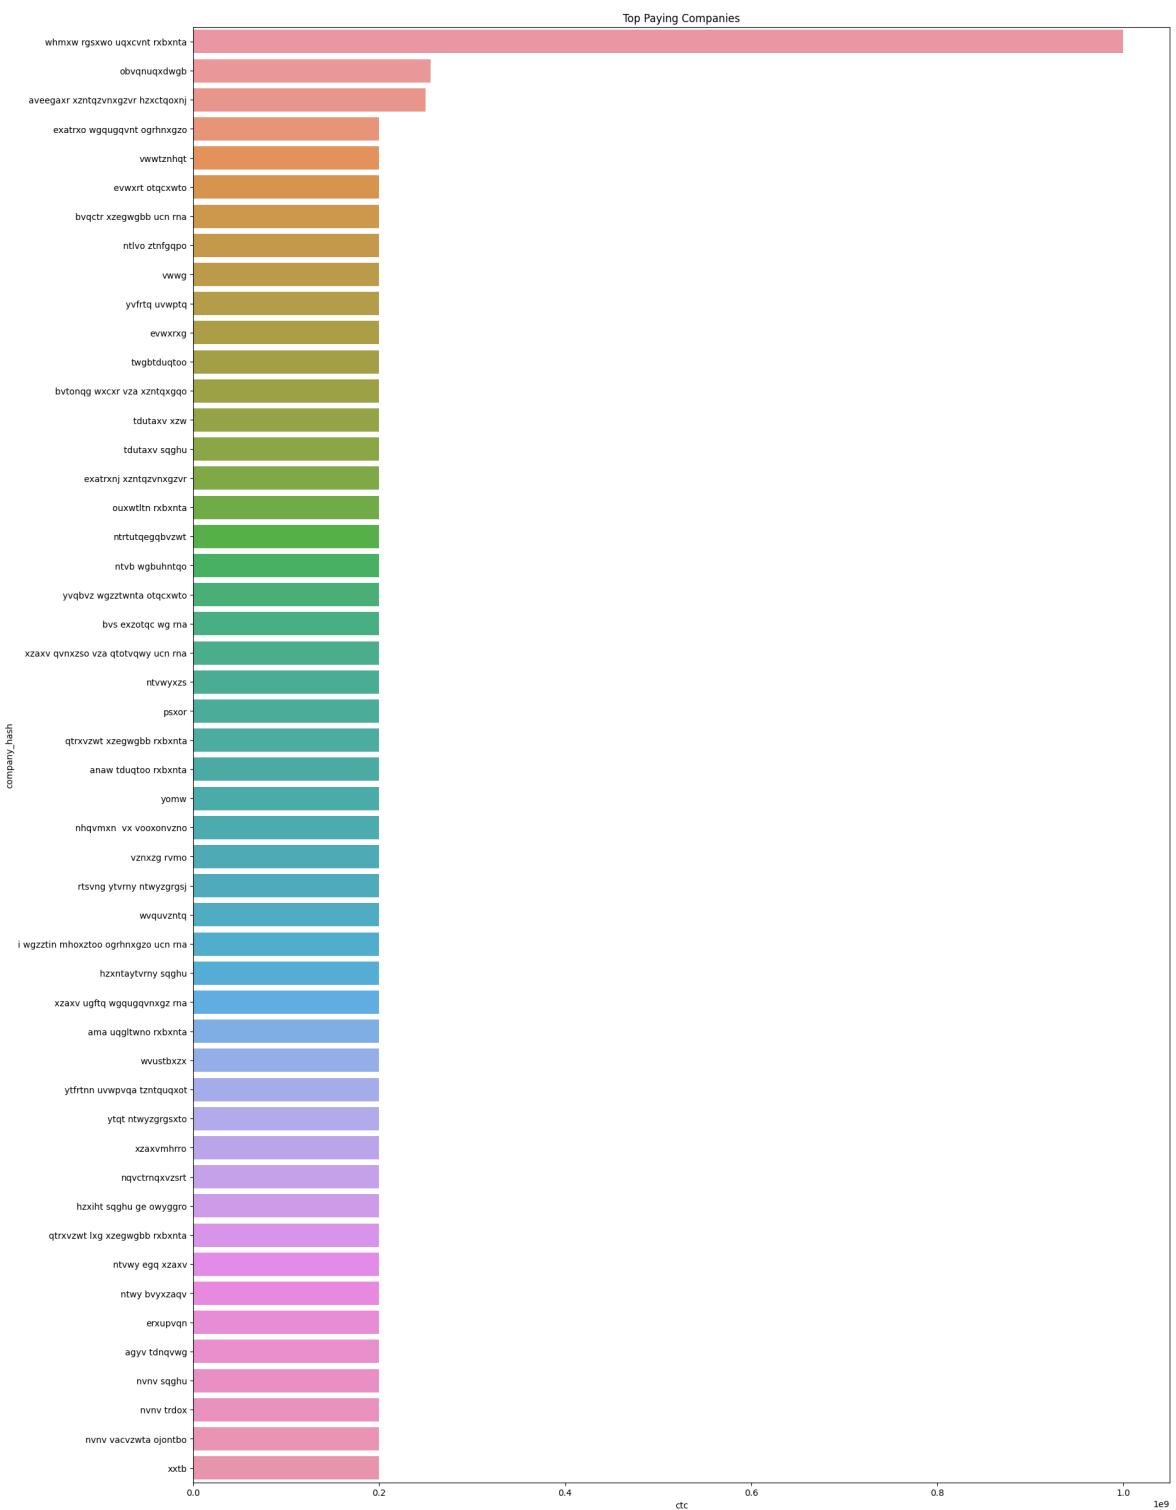
```
In [33]:  
tmp = df.copy()  
tmp = tmp.groupby(['job_position']).max()['ctc'].reset_index().sort_values('ctc')  
plt.figure(figsize=(20,30))  
sns.barplot(data=tmp,x='ctc',y='job_position').set(title="Top Paying Jobs")  
plt.show()
```



```
In [34]: list(tmp['job_position'])
```

```
Out[34]: ['nan',
'database administrator',
'backend engineer',
'other',
'sdet',
'data analyst',
'data scientist',
'qa engineer',
'support engineer',
'non coder',
'devops engineer',
'engineering intern',
'cofounder',
'engineering leadership',
'android engineer',
'security leadership',
'fullstack engineer',
'program manager',
'ios engineer',
'product manager',
'backend architect',
'product designer',
'frontend engineer',
'research engineers',
'release engineer',
'driver',
'business man',
'student',
'reseller',
'assistant manager',
'data entry',
'jharkhand',
'telar',
'owner',
'teaching',
'computer scientist',
'safety officer',
'seleceman',
'senior system engineer',
'software test engineer',
'computer faculty',
'applications engineer',
'sde',
'security intern',
'electric power supply',
'intern software developer',
'toyota',
'software enginner',
'veice president',
'phd student']
```

```
In [35]: tmp = df.copy()
tmp = tmp.groupby(['company_hash']).max()['ctc'].reset_index().sort_values('ctc')
plt.figure(figsize=(20,30))
sns.barplot(data=tmp,x='ctc',y='company_hash').set(title="Top Paying Companies")
plt.show()
```



```
In [36]: list(tmp['company_hash'])
```

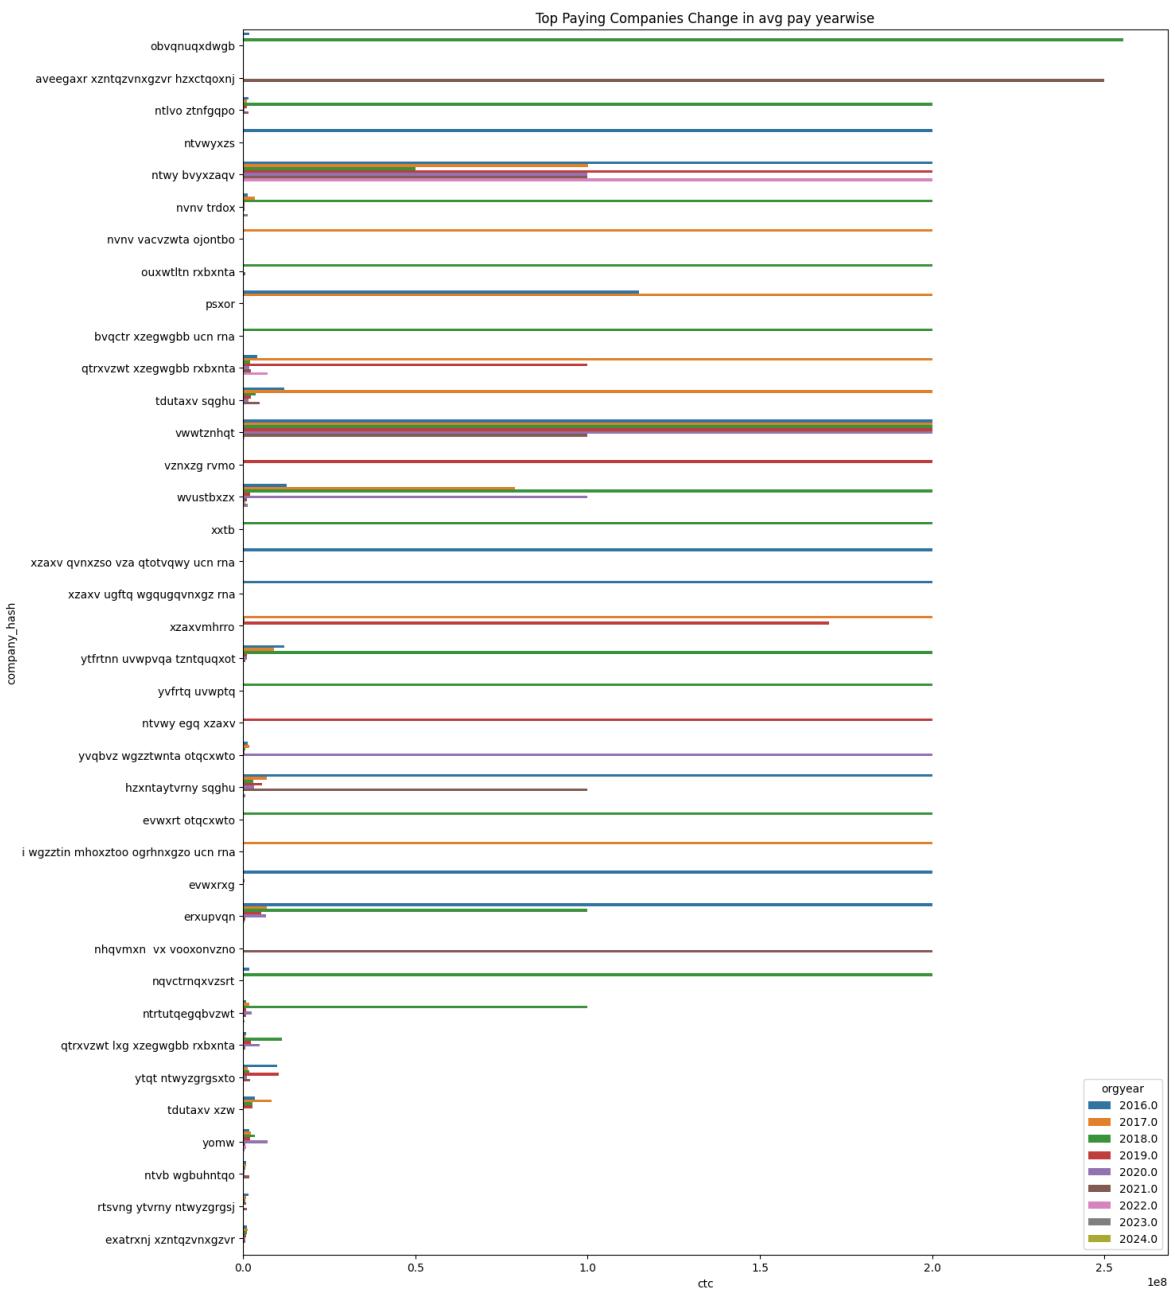
```
Out[36]: ['whmxw rgsxwo uqxcvnt rxbxnta',
 'obvqnuqxdwgb',
 'aveegaxr xzntqzvnxgzvr hzxctqoxnj',
 'exatrxo wgqugqvnt ogrhnxgzo',
 'vwwtznhqt',
 'evwxrt otqcxwto',
 'bvqctr xzegwgb ucn rna',
 'ntlvo ztnfgqpo',
 'vwwg',
 'yvfrtq uvwptq',
 'evwxrxg',
 'twgbtduqtoo',
 'bvtionqg wxcxr vza xzntqxgqo',
 'tdutaxv xzw',
 'tdutaxv sqghu',
 'exatrxnj xzntqzvnxgzvr',
 'ouxwtln rxbxnta',
 'ntrtutqegqbvzwt',
 'ntvb wgbuhntqo',
 'yvqbvz wgzztwnta otqcxwto',
 'bvs exzotqc wg rna',
 'xzaxv qvnxzso vza qtotvqwy ucn rna',
 'ntvwyxzs',
 'psxor',
 'qtrxvzwt xzegwgb rxbxnta',
 'anaw tduqtoo rxbxnta',
 'yomw',
 'nhqvmxn vx vooxonvzno',
 'vznxzg rvmo',
 'rtsvng ytvrny ntwyzgrgsj',
 'wvquvzntq',
 'i wgzztin mhoxztoo ogrhnxgzo ucn rna',
 'hzxntaytvrny sqghu',
 'xzaxv ugftq wgqugqvnxgz rna',
 'ama uqgltwo rxbxnta',
 'wvustbxzx',
 'ytfrttn uvwpvqa tzntquqxot',
 'ytqt ntwyzgrgsxto',
 'xzaxvmlhrro',
 'nqvctrnqvxzsrt',
 'hzxiht sqghu ge owyggro',
 'qtrxvzwt lxr xzegwgb rxbxnta',
 'ntvwy egg xzaxv',
 'ntwy bvyxzaqv',
 'erxupvqn',
 'agyv tdnqvwg',
 'nnnv sqghu',
 'nnnv trdox',
 'nnnv vacvzwta ojontbo',
 'xxtb']
```

```
In [37]: tmp = df.copy()
tmp = tmp[tmp['company_hash'].isin(['whmxw rgsxwo uqxcvnt rxbxnta',
 'obvqnuqxdwgb',
 'aveegaxr xzntqzvnxgzvr hzxctqoxnj',
 'exatrxo wgqugqvnt ogrhnxgzo',
 'vwwtznhqt',
 'evwxrt otqcxwto',
 'bvqctr xzegwgb ucn rna',
 'ntlvo ztnfgqpo',
```

```

'vwwg',
'yvfrtq uvwptq',
'evwxrxg',
'twgbtduqtoo',
'bvtong wxcxr vza xzntqgxgqo',
'tdutaxv xzw',
'tdutaxv sqghu',
'exatrxnj xzntqzvnxgzvr',
'ouxwtltn rxbxnta',
'ntrtutqeqqbvzwt',
'ntvb wgbuhntqo',
'yvqbvz wgzztwnta otqcxwto',
'bvs exzotqc wg rna',
'xzaxv qvnxzso vza qtotvqwy ucn rna',
'ntvwyxzs',
'psxor',
'qtrxvzwt xzegwgb rxbxnta',
'anaw tduqtoo rxbxnta',
'yomw',
'nhqvmxn vx vooxonvzno',
'vznxzg rvmo',
'rtsvng ytvrny ntwyzgrgsj',
'wvquvzntq',
'i wgzztin mhoxztoo ogrhnxgzo ucn rna',
'hzxntaytvnry sqghu',
'xzaxv ugftq wgqugqvnxgz rna',
'ama uqgltwo rxbxnta',
'wvustbxzx',
'ytfrtnn uvwpvqa tzntquqxot',
'ytqt ntwyzgrgsxto',
'xzaxvmlhrro',
'nqvctrnqvxzsrt',
'hzxiht sqghu ge owyggro',
'qtrxvzwt lxr xzegwgb rxbxnta',
'ntvwy egg xzaxv',
'ntwy bvyxzaqv',
'erxupvqn',
'agyv tdnqvwg',
'nvnv sqghu',
'nvnv trdox',
'nvnv vacvzwta ojontbo',
'xxtb'])]
tmp = tmp[tmp['orgyear'] >= 2016]
tmp = tmp.groupby(['company_hash', 'orgyear']).max()['ctc'].reset_index().sort_values(by='ctc', ascending=False)
plt.figure(figsize=(15, 20))
sns.barplot(data=tmp, x='ctc', y='company_hash', hue='orgyear').set(title="Top Paying Companies by CTC in 2016")
plt.show()

```

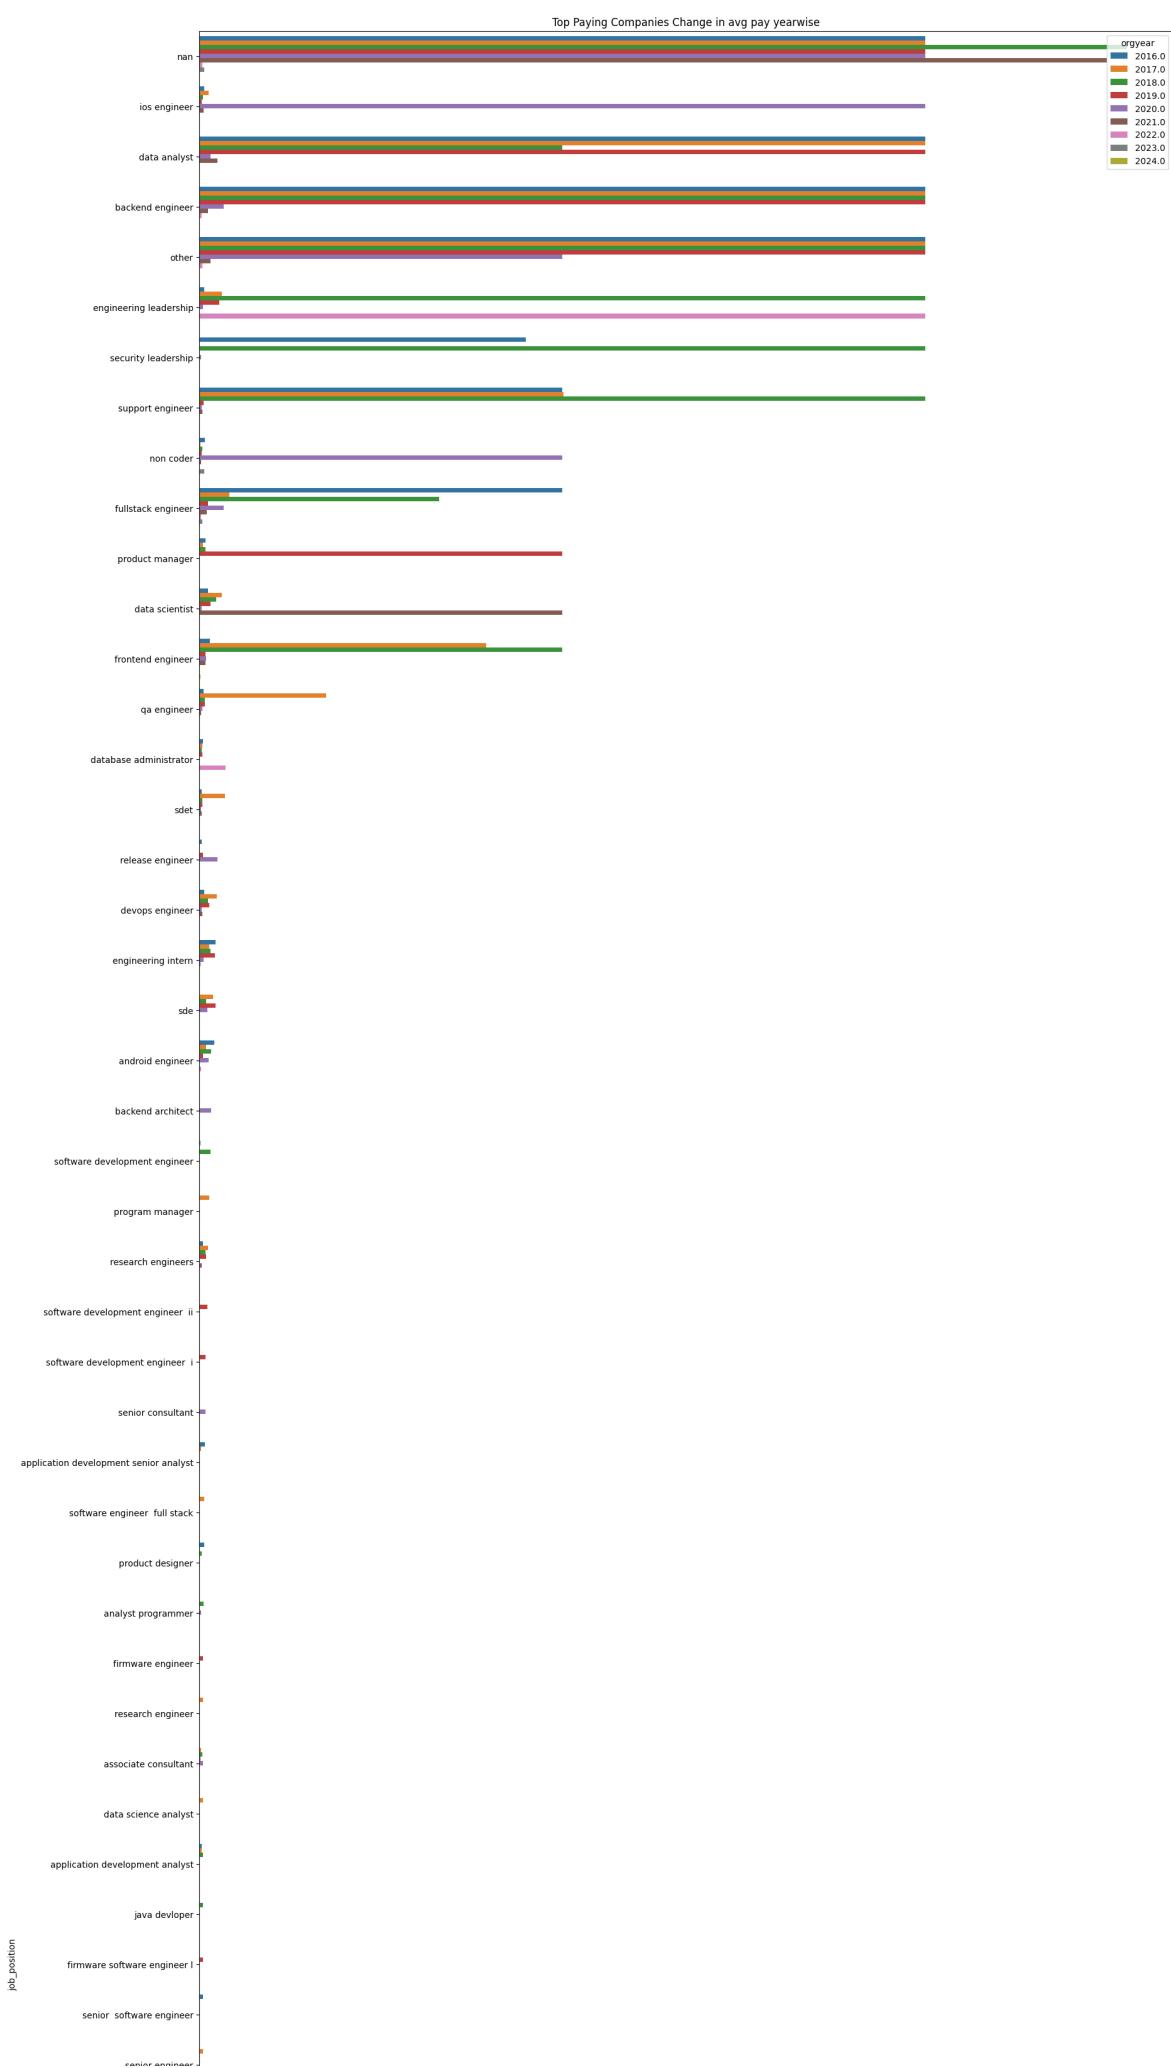


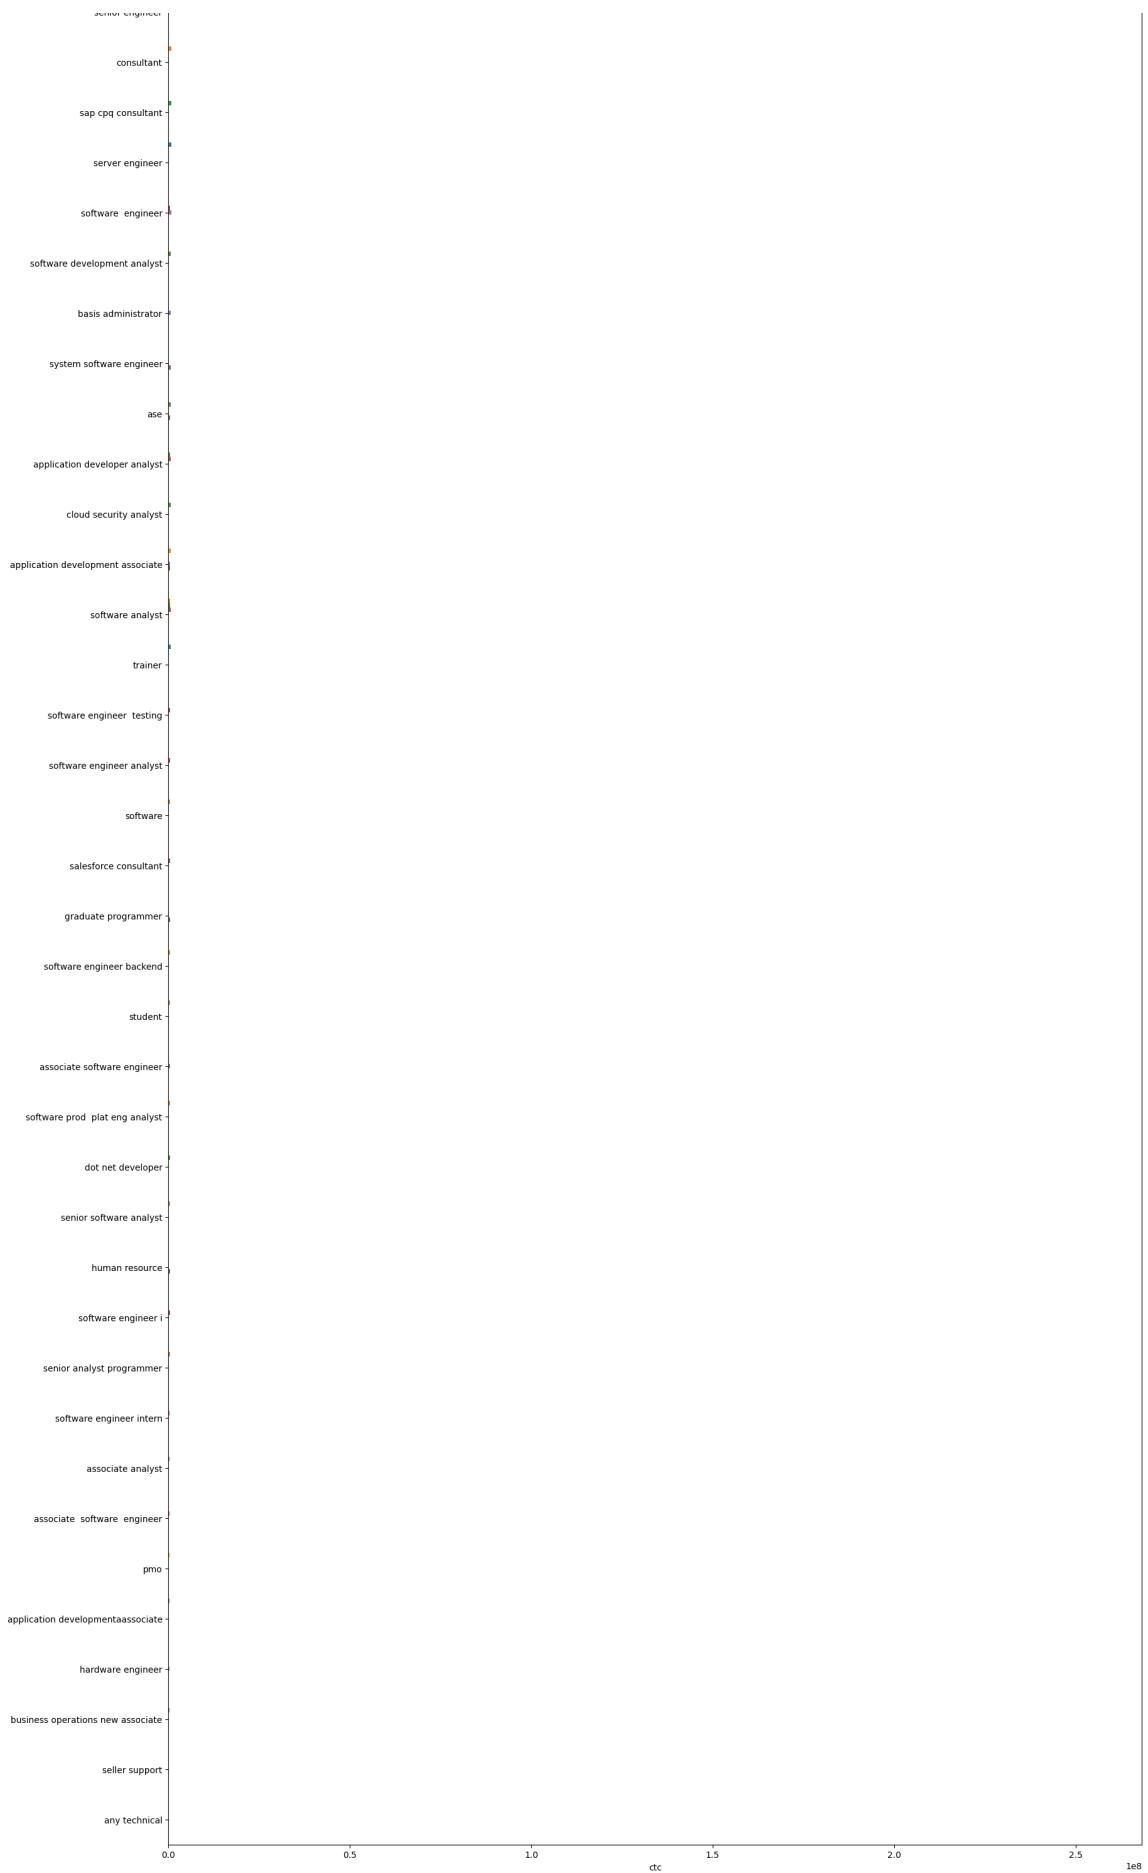
```
In [38]: tmp = df.copy()
tmp = tmp[tmp['company_hash'].isin(['whmxw rgsxwo uqxcvnt rxbxnta',
 'obvqnuqxdwgb',
 'aveegaxr xzntqzvnxgvr hzxctqoxnj',
 'exatrxo wgqugqvnt ogrhnxgzo',
 'vwwtzhqht',
 'evwxrt otqcxwto',
 'bvqctr xzegwgbb ucn rna',
 'ntlvo ztnfgqpo',
 'vwwg',
 'yvfrtq uwptq',
 'evwxrxg',
 'twgbtduqtoo',
 'bvtong wxcxr vza xzntqzgqo',
 'tdutaxv xzw',
 'tdutaxv sqghu',
 'exatrxnj xzntqzvnxgvr',
 'ouxwtln rxbxnta',
 'ntrtutqeqqbzwt',
 'ntvb wgbuhntqo',
 'yvqbzv wgzztwnta otqcxwto',
```

```

'bvs exzotqc wg rna',
'xzaxv qvnxzso vza qtovqwy ucn rna',
'ntvwyxzs',
'psxor',
'qtrxzwt xzegwgbp rxbxnta',
'anaw tduqtoo rxbxnta',
'yomw',
'nhqvmxn vx vooxonvzno',
'vznxzg rvmo',
'rtsvng ytvrny ntwygrgsj',
'wvquvzntq',
'i wgzztin mhoxztoo ogrhnxgzo ucn rna',
'hzxntaytvnry sqghu',
'xzaxv ugftq wgquugqvnxgz rna',
'ama uqgltwno rxbxnta',
'wvustbxzx',
'ytfrtnn uvwpvqa tzntquqxot',
'ytqt ntwygrgsxto',
'xzaxvvhro',
'nqvctrnqvxzsrt',
'hzxiht sqghu ge owyggro',
'qtrxzwt lxr xzegwgbp rxbxnta',
'ntvwy egg xzaxv',
'ntwy bvyxzaqv',
'erxupvqn',
'agyv tdnqvwg',
'nvnv sqghu',
'nvnv trdox',
'nvnv vacvzwta ojontbo',
'xxtb'])]
tmp = tmp[tmp['orgyear'] >= 2016]
tmp = tmp.groupby(['job_position','orgyear']).max()['ctc'].reset_index().sort_values(by='ctc', ascending=False)
plt.figure(figsize=(20,80))
sns.barplot(data=tmp,x='ctc',y='job_position',hue='orgyear').set(title="Top Paying Job Positions by CTC")
plt.show()

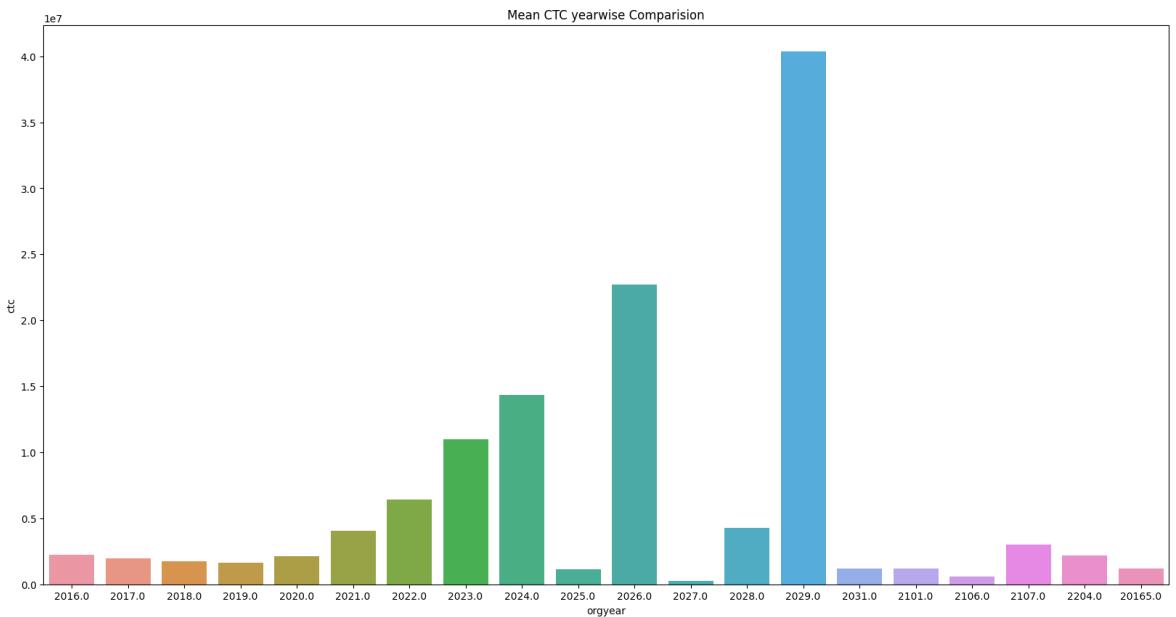
```





```
In [39]: tmp = df.copy()
tmp = tmp[tmp['orgyear'] >= 2016]
tmp = tmp.groupby(['orgyear']).mean()['ctc'].reset_index().sort_values('ctc', ascending=True)
plt.figure(figsize=(20,10))
```

```
sns.barplot(data=tmp,y='ctc',x='orgyear').set(title="Mean CTC yearwise Comparision")
plt.show()
```



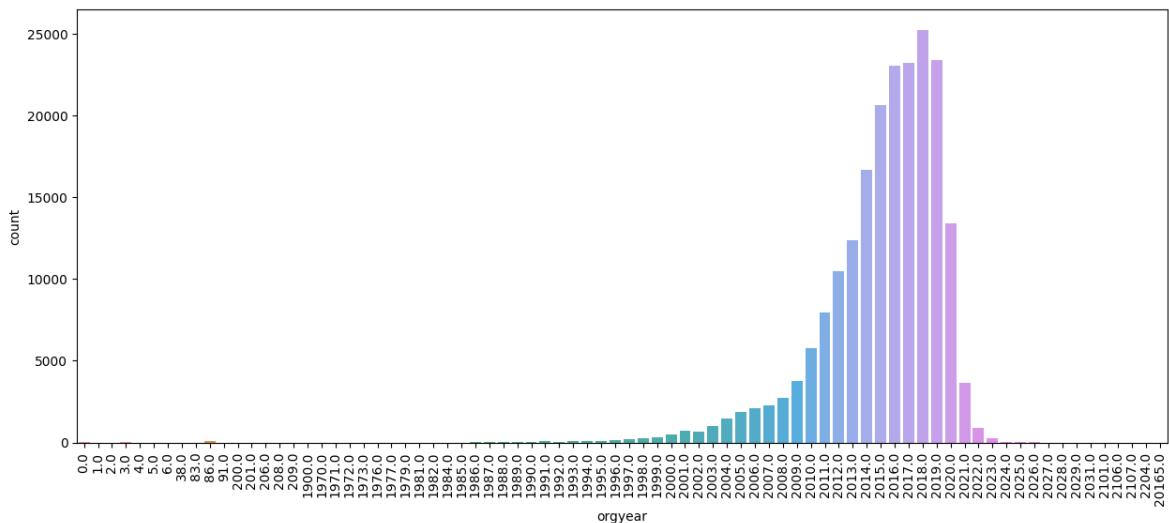
Outlier Treatment

Orgyear

```
In [40]: df["orgyear"].value_counts()
```

```
Out[40]: 2018.0    25240
2019.0    23402
2017.0    23237
2016.0    23038
2015.0    20602
...
2107.0      1
1972.0      1
2101.0      1
208.0       1
200.0       1
Name: orgyear, Length: 78, dtype: int64
```

```
In [41]: plt.figure(figsize=(15,6))
sns.countplot(x=df['orgyear'])
plt.xticks(rotation = 90)
plt.show()
```



```
In [42]: df["orgyear"].quantile(0.001)
```

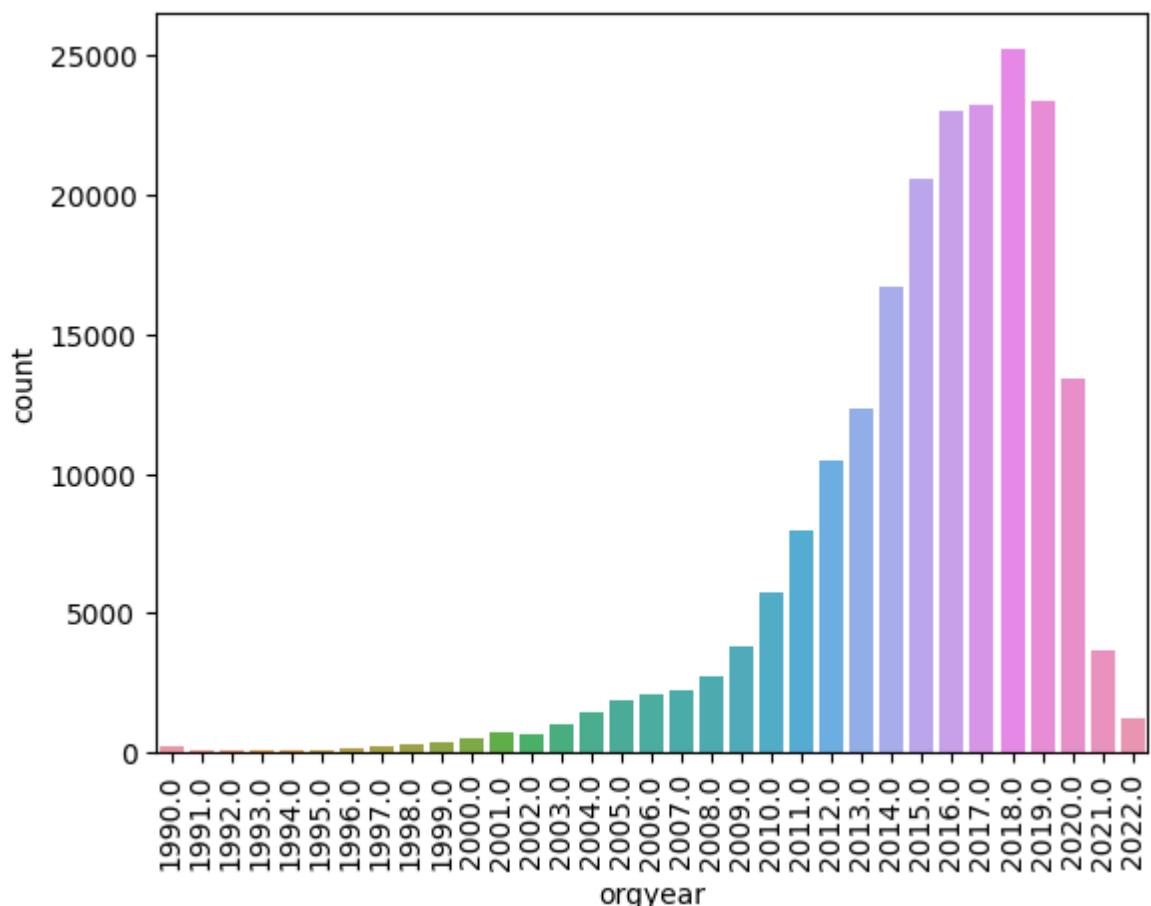
```
Out[42]: 1990.0
```

```
In [43]: df["orgyear"].quantile(0.999)
```

```
Out[43]: 2023.0
```

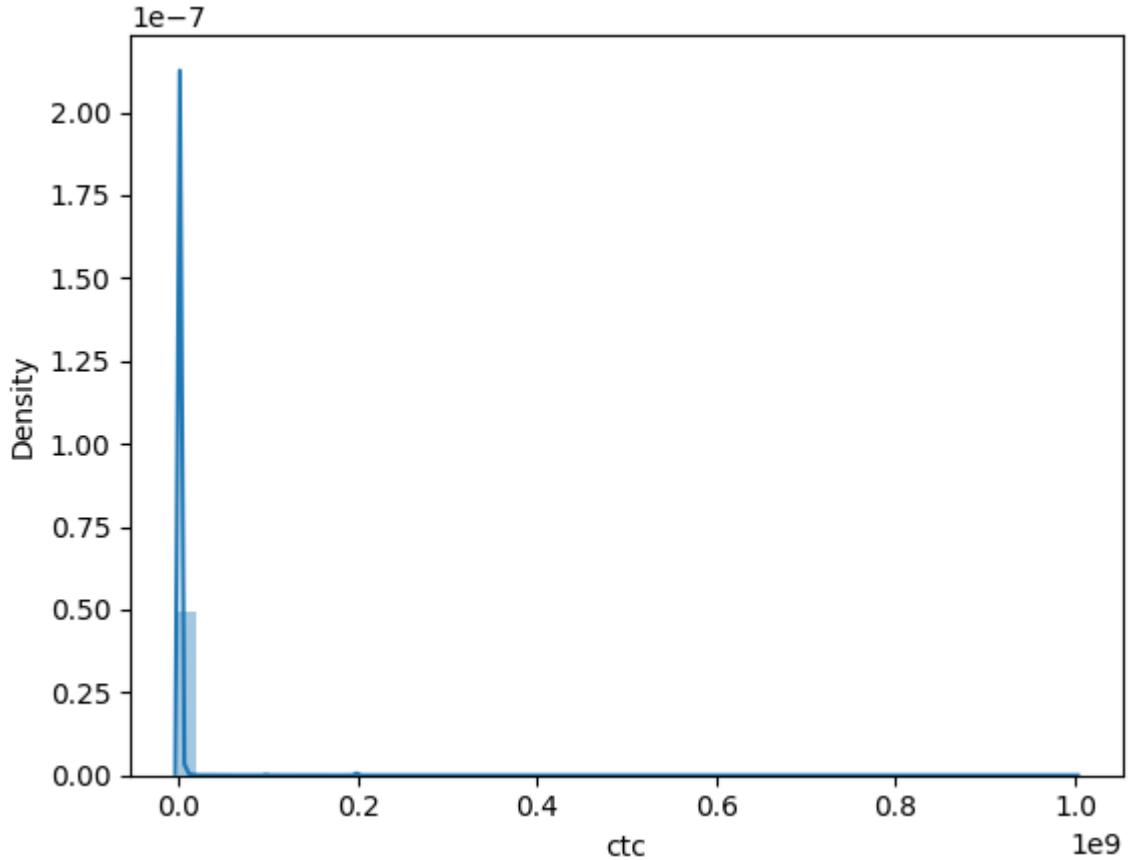
```
In [44]: df["orgyear"] = df["orgyear"].clip(1990, 2022)
```

```
In [45]: sns.countplot(x=df['orgyear'])
plt.xticks(rotation = 90)
plt.show()
```

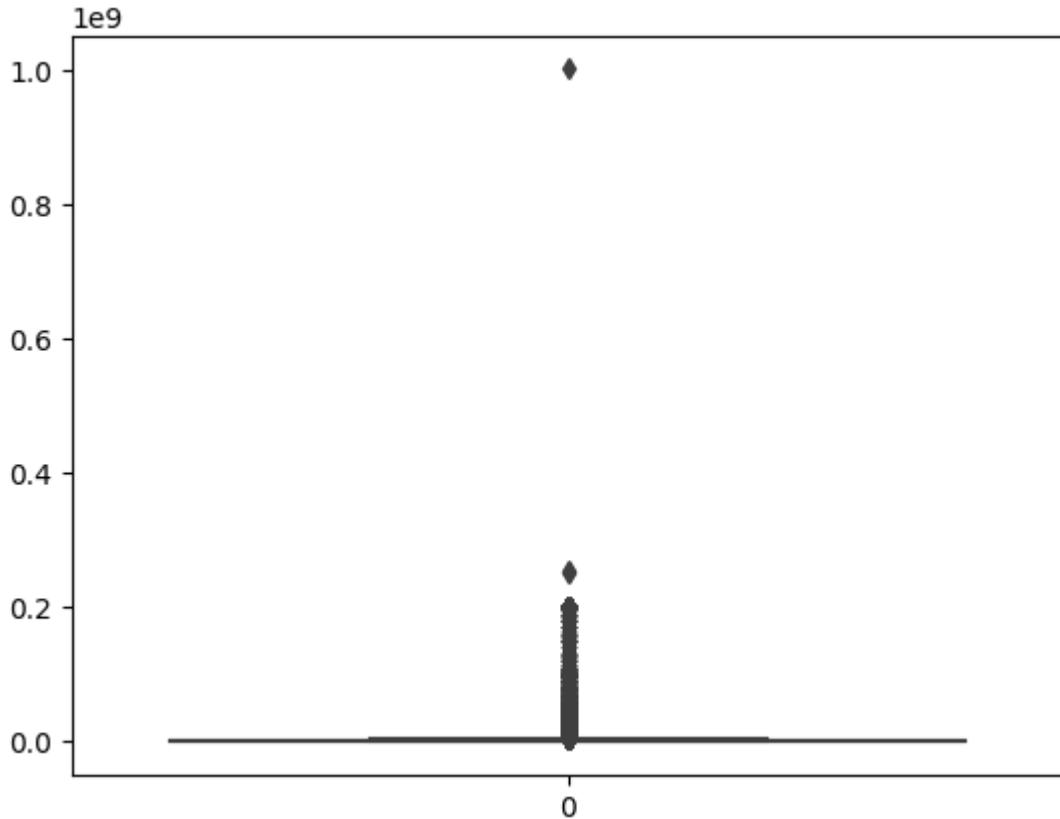


CTC

```
In [46]: sns.distplot(df["ctc"])
plt.show()
```



```
In [47]: sns.boxplot(df["ctc"])
plt.show()
```



```
In [48]: df["ctc"].quantile(0.01)
```

```
Out[48]: 37000.0
```

```
In [49]: df["ctc"].quantile(0.999)
```

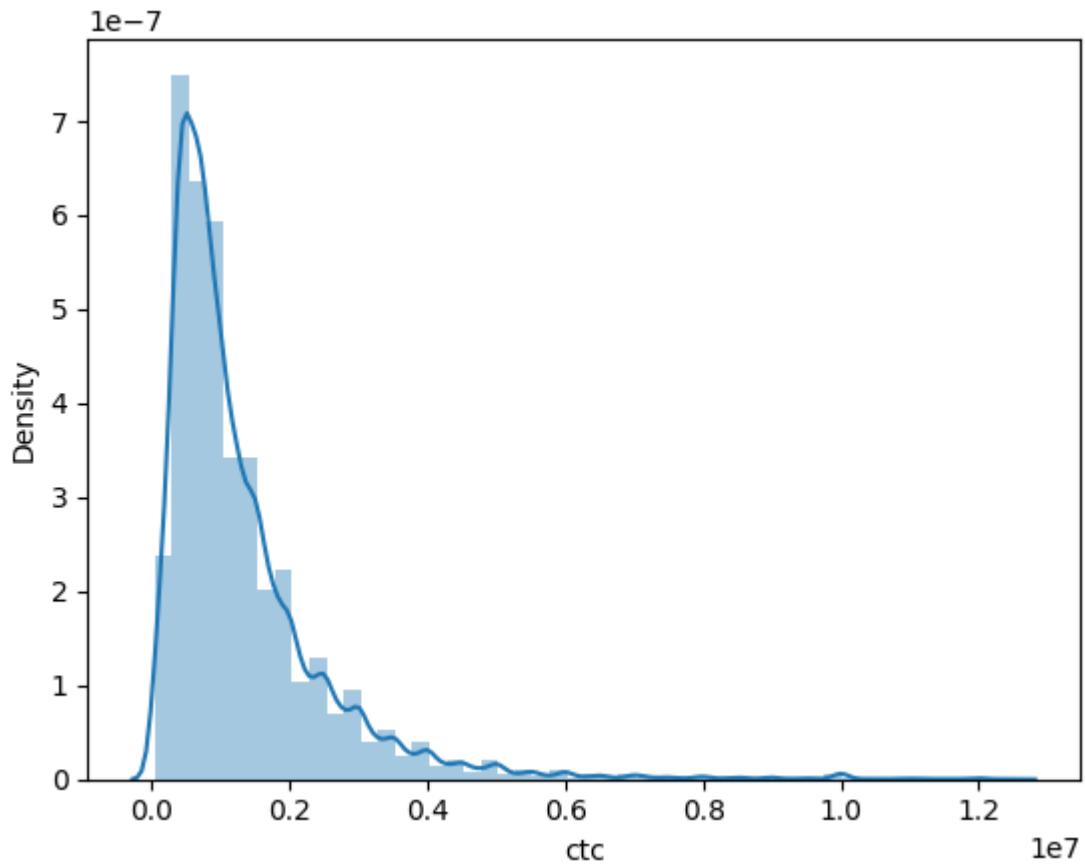
```
Out[49]: 200000000.0
```

```
In [50]: df = df.loc[((df.ctc) > df.ctc.quantile(0.01)) & ((df.ctc) < df.ctc.quantile(0.999))]
```

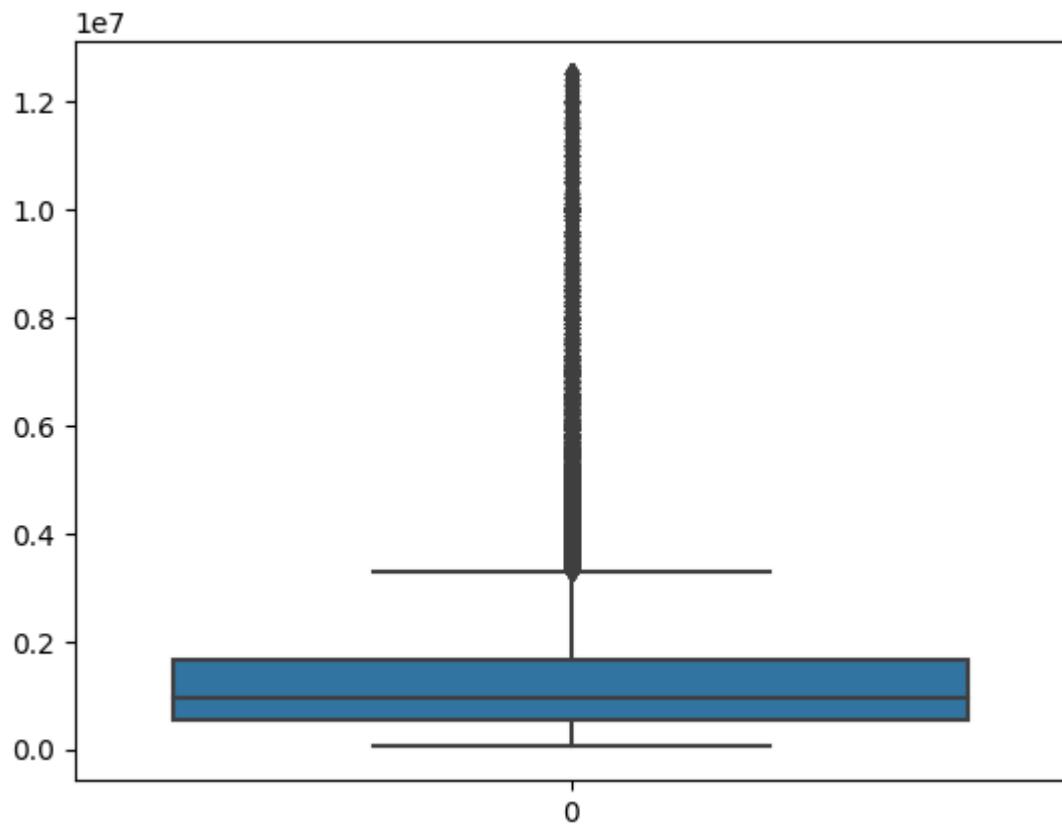
```
In [51]: df.head()
```

	company_hash	email_hash	orgyear	ctc
0	atrgxnnt xzaxv	6de0a4417d18ab14334c3f43397fc13b30c35149d70c05...	2016.0	1100000
1	qtrxzvwt xzegwgbbr rxbxnta	b0aa1ac138b53cb6e039ba2c3d6604a250d02d5145c10...	2018.0	449999
2	ojzwnvwnxw vx	4860c670bcd48fb96c02a4b0ae3608ae6fdd98176112e9...	2015.0	2000000
3	ngpgutaxv	effdede7a2e7c2af664c8a31d9346385016128d66bbc58...	2017.0	700000
4	qxen sqghu	6ff54e709262f55cb999a1c1db8436cb2055d8f79ab520...	2017.0	1400000

```
In [52]: sns.distplot(df["ctc"])
plt.show()
```



```
In [53]: sns.boxplot(df["ctc"])
plt.show()
```



ctc updated_year

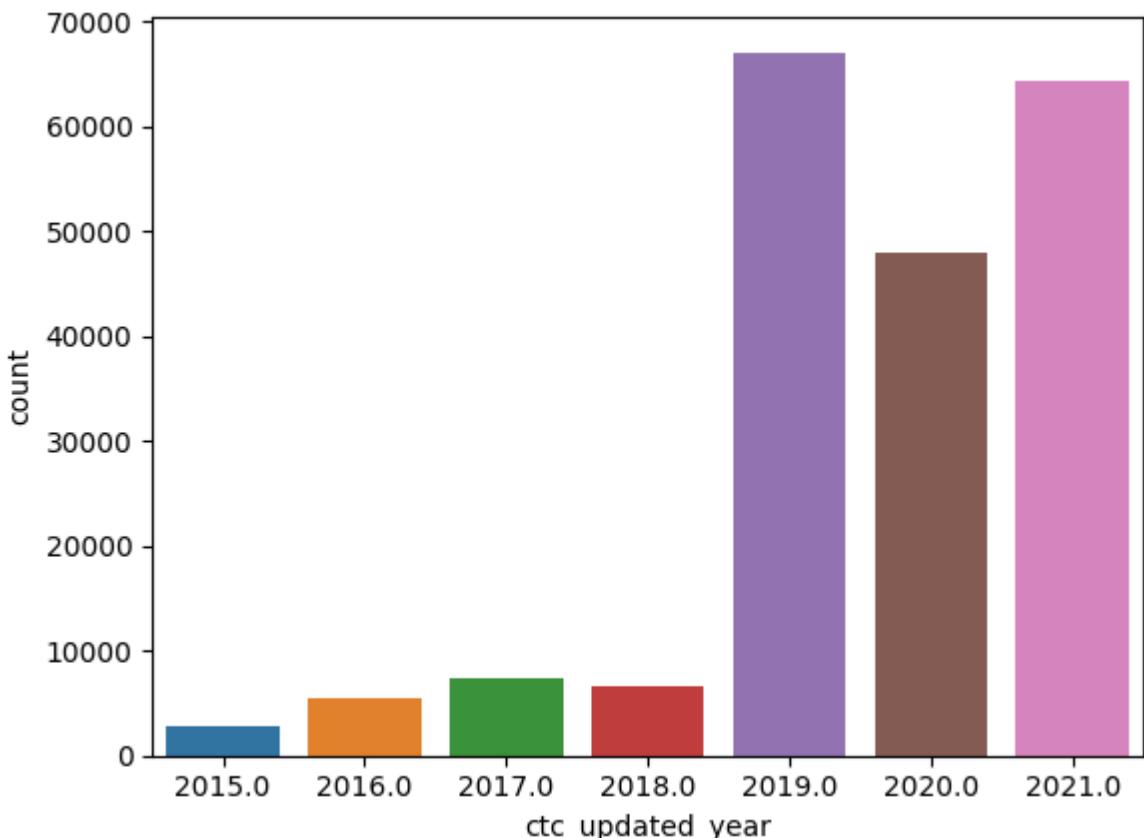
```
In [54]: df["ctc_updated_year"].quantile(0.001)
```

```
Out[54]: 2015.0
```

```
In [55]: df["ctc_updated_year"].quantile(0.99)
```

```
Out[55]: 2021.0
```

```
In [56]: sns.countplot(x=df['ctc_updated_year'])
plt.show()
```



```
In [57]: df.loc[df['job_position']=='nan', 'job_position']=np.nan
```

```
In [58]: df.loc[df["company_hash"]=="nan","company_hash"] = np.nan
```

Feature Engineering

Masked company name to "Others" having count less than 5

```
In [59]: df.loc[df.groupby("company_hash")["ctc"].transform("count") < 5,"company_hash"]
```

```
In [60]: (df["company_hash"] == "Others").sum()
```

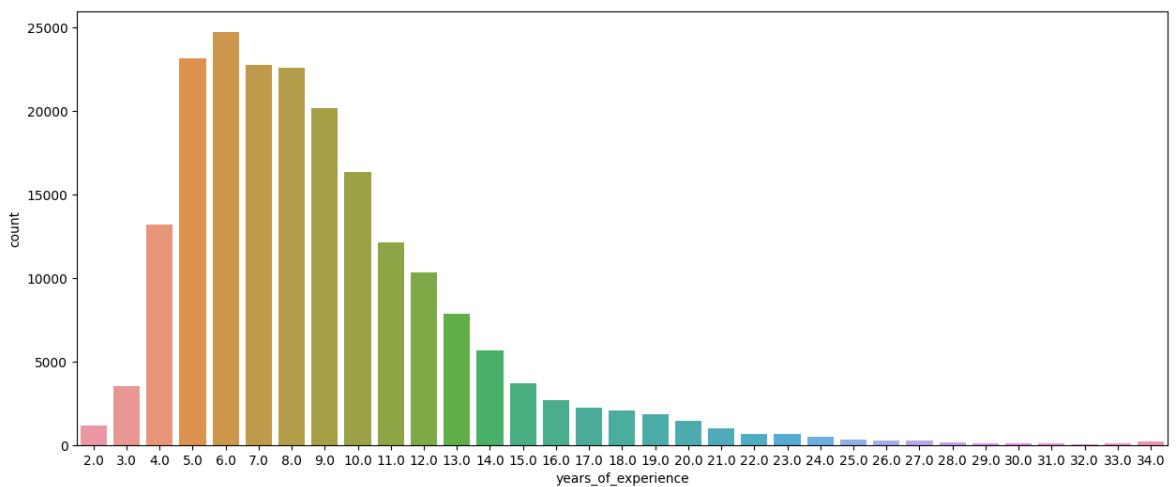
```
Out[60]: 46434
```

```
In [61]: df['orgyear'].describe()
```

```
Out[61]: count    201625.000000
          mean     2015.104769
          std      4.256063
          min     1990.000000
          25%    2013.000000
          50%    2016.000000
          75%    2018.000000
          max     2022.000000
          Name: orgyear, dtype: float64
```

```
In [62]: # years of experience = current year - employment start year
df["years_of_experience"] = 2024 - df["orgyear"]
```

```
In [63]: plt.figure(figsize=(15,6))
sns.countplot(x=df['years_of_experience'])
plt.show()
```



```
In [64]: df.duplicated().sum()
```

```
Out[64]: 212
```

```
In [65]: df.drop_duplicates(inplace=True)
```

```
In [66]: df.isna().sum()
```

```
Out[66]: company_hash           42
          email_hash            0
          orgyear                0
          ctc                     0
          job_position          51671
          ctc_updated_year        0
          years_of_experience      0
          dtype: int64
```

```
In [67]: # records having ctc_updated_year higher than their organization joining year
(df["ctc_updated_year"] < df["orgyear"]).sum()
```

```
Out[67]: 8465
```

```
In [68]: df.ctc_updated_year = df[["ctc_updated_year", "orgyear"]].max(axis = 1)
```

```
In [69]: (df["ctc_updated_year"] < df["orgyear"]).sum()
```

```
Out[69]: 0
```

```
In [70]: df.sample(2)
```

	company_hash		email_hash	orgyear	
11711	Others	feed3529e8360ae24b721dec3eac9124b104886d36b722...		2018.0	
109030	ovbohzs qa xzonxnhnt xzaxv atryx	fc1ffb7be86e45f41c95b140b066e2a639ac54d10cf096...		2018.0	1

```
In [71]: # Filling null values with others -- if not done before  
df['job_position'] = df['job_position'].fillna('Others')  
df['company_hash'] = df['company_hash'].fillna('Others')
```

```
In [72]: df.isna().sum()
```

```
Out[72]: company_hash      0  
email_hash        0  
orgyear          0  
ctc              0  
job_position     0  
ctc_updated_year 0  
years_of_experience 0  
dtype: int64
```

```
In [73]: df.duplicated().sum()
```

```
Out[73]: 103
```

```
In [74]: df.describe()
```

```
Out[74]:
```

	orgyear	ctc	ctc_updated_year	years_of_experience
count	201413.000000	2.014130e+05	201413.000000	201413.000000
mean	2015.103722	1.313212e+06	2019.687314	8.896278
std	4.257499	1.234111e+06	1.287119	4.257499
min	1990.000000	3.800000e+04	2015.000000	2.000000
25%	2013.000000	5.500000e+05	2019.000000	6.000000
50%	2016.000000	9.500000e+05	2020.000000	8.000000
75%	2018.000000	1.650000e+06	2021.000000	11.000000
max	2022.000000	1.250000e+07	2022.000000	34.000000

```
In [75]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 201413 entries, 0 to 205842
Data columns (total 7 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   company_hash     201413 non-null   object  
 1   email_hash       201413 non-null   object  
 2   orgyear          201413 non-null   float64 
 3   ctc              201413 non-null   int64  
 4   job_position     201413 non-null   object  
 5   ctc_updated_year 201413 non-null   float64 
 6   years_of_experience 201413 non-null   float64 
dtypes: float64(3), int64(1), object(3)
memory usage: 12.3+ MB

```

Manual Clustering

```
In [76]: grp = ['company_hash', 'job_position', 'years_of_experience']
data_tmp1 = df.groupby(grp).agg({'ctc':[ 'mean', 'median', 'min', 'max', 'count' ]})
data_tmp1.columns = ["{} {}".format(b_, a_) if a_ not in grp else "{}".format(a) for b_ in data_tmp1.columns for a_ in data_tmp1.columns]
data_tmp1.head()
```

Out[76]:

	company_hash	job_position	years_of_experience	mean ctc	median ctc	min ctc	n	
0	Others	Others		2.0	1.492754e+06	800000.0	60000	100
1	Others	Others		3.0	9.209444e+05	650000.0	47000	100
2	Others	Others		4.0	8.635434e+05	550000.0	40000	100
3	Others	Others		5.0	7.519201e+05	500000.0	40000	100
4	Others	Others		6.0	6.756606e+05	500000.0	40000	100

```
In [77]: datatmp = df.merge(data_tmp1[['company_hash', 'job_position', 'years_of_experience']], left_index=True, right_index=True)
col1 = 'ctc'
col2 = 'mean ctc'
conditions = [ datatmp[col1] > datatmp[col2], datatmp[col1] == datatmp[col2], datatmp[col1] < datatmp[col2] ]
choices = [ 1, 2, 3 ]

datatmp['Designation'] = np.select(conditions, choices, default=np.nan)
datatmp.head()
```

Out[77]:

	company_hash		email_hash	orgyear	ctc
0	atrgxnnt xzaxv	6de0a4417d18ab14334c3f43397fc13b30c35149d70c05...		2016.0	1100000
1	qtrxvzwt xzegwgbbr rxbxnta	b0aaf1ac138b53cb6e039ba2c3d6604a250d02d5145c10...		2018.0	449999
2	Others	4860c670bcd48fb96c02a4b0ae3608ae6fdd98176112e9...		2015.0	2000000
3	ngpgutaxv	effdede7a2e7c2af664c8a31d9346385016128d66bbc58...		2017.0	700000
4	qxen sqghu	6ff54e709262f55cb999a1c1db8436cb2055d8f79ab520...		2017.0	1400000

In [78]:

```
# unique value designation column(listed in %)
designation = datatmp['Designation'].value_counts(normalize=True).map(lambda cal
designation.columns = ['Designation', 'Count']
designation
```

Out[78]:

	Designation	Count
0	3.0	46.16
1	1.0	30.33
2	2.0	23.52

Manual Clustering on company_hash and job position

In [79]:

```
grp = ['company_hash', 'job_position']
data_tmp1 = datatmp.groupby(grp).agg({'ctc':[('mean2', 'mean'), 'median', 'min', 'ma
data_tmp1.columns  = ["{} {}".format(b_, a_) if a_ not in grp else "{}".format(a
data_tmp1.head()
```

Out[79]:

	company_hash	job_position	mean2 ctc	median ctc	min ctc	max ctc	count ctc
0	Others	Others	9.977383e+05	700000.0	40000	12500000	9450
1	Others	a group chat application	5.000000e+05	500000.0	500000	500000	1
2	Others	abap developer	5.000000e+05	500000.0	500000	500000	1
3	Others	administrative clerk	5.000000e+05	500000.0	500000	500000	1
4	Others	administrator	1.940000e+06	1940000.0	380000	3500000	2

In [80]:

```
datatmp = datatmp.merge(data_tmp1[grp + ['mean2 ctc']], on=grp, how='left')
col1 = 'ctc'
col2 = 'mean2 ctc'
conditions = [ datatmp[col1] > datatmp[col2], datatmp[col1] == datatmp[col2], d
```

```

choices      = [ 1, 2, 3 ]

datatmp['Class'] = np.select(conditions, choices, default=np.nan)
datatmp.head()

```

Out[80]:

	company_hash	email_hash	orgyear	ctc
0	atrgxnnt xzaxv	6de0a4417d18ab14334c3f43397fc13b30c35149d70c05...	2016.0	1100000
1	qtrxvzwt xzegwgbb rxbxnta	b0aaaf1ac138b53cb6e039ba2c3d6604a250d02d5145c10...	2018.0	449999
2	Others	4860c670bcd48fb96c02a4b0ae3608ae6fdd98176112e9...	2015.0	2000000
3	npgputaxv	effdede7a2e7c2af664c8a31d9346385016128d66bbc58...	2017.0	700000
4	qxen sqghu	6ff54e709262f55cb999a1c1db8436cb2055d8f79ab520...	2017.0	1400000

In [81]:

```

# unique value Class column(Listed in %)
Class = datatmp['Class'].value_counts(normalize=True).map(lambda calc: round(100*calc, 2))
Class.columns = ['Class', 'Count']
Class

```

Out[81]:

	Class	Count
0	3.0	57.98
1	1.0	35.78
2	2.0	6.24

Manual Clustering based on company

based on ctc per company , assigning company as tier 1 2 and 3 per each learners

In [82]:

```

grp = ['company_hash']
data_tmp1 = datatmp.groupby(grp).agg({'ctc':[('mean3','mean'), 'median', 'min', 'max']})
data_tmp1.columns = ["{} {}".format(b_, a_) if a_ not in grp else "{}".format(a) for b_, a_ in data_tmp1.columns]
data_tmp1.head()

```

Out[82]:

	company_hash	mean3 ctc	median ctc	min ctc	max ctc	count ctc
0	Others	1.108693e+06	770000.0	38000	12500000	46376
1	a ntwygrgsxto	1.234688e+06	600000.0	350000	4000000	16
2	aaqxctz avnv owxtzwto vzvrjnxwo ucn rna	9.850000e+05	500000.0	360000	3600000	8
3	abwavnv ojontb	7.320000e+05	700000.0	700000	780000	5
4	adw ntwygrgsj	9.098081e+05	600000.0	56000	8000000	297

In [83]:

```
datatmp = datatmp.merge(data_tmp1[grp + ['mean3 ctc']], on=grp, how='left')
col1 = 'ctc'
col2 = 'mean3 ctc'
conditions = [ datatmp[col1] > datatmp[col2], datatmp[col1] == datatmp[col2], d
choices = [ 1, 2, 3 ]

datatmp['Tier'] = np.select(conditions, choices, default=np.nan)
datatmp.head()
```

Out[83]:

	company_hash	email_hash	orgyear	ctc
0	atrgxnnt xzaxv	6de0a4417d18ab14334c3f43397fc13b30c35149d70c05...	2016.0	1100000
1	qtrxvzwt xzegwgbb rxbxnta	b0aaf1ac138b53cb6e039ba2c3d6604a250d02d5145c10...	2018.0	449999
2	Others	4860c670bcd48fb96c02a4b0ae3608ae6fd98176112e9...	2015.0	2000000
3	ngpgutaxv	effdede7a2e7c2af664c8a31d9346385016128d66bbc58...	2017.0	700000
4	qxen sqghu	6ff54e709262f55cb999a1c1db8436cb2055d8f79ab520...	2017.0	1400000

In [84]:

```
# unique value Tier column(listed in %)
Tier = datatmp['Tier'].value_counts(normalize=True).map(lambda calc: round(100*calc))
Tier.columns = ['Tier', 'Count']
Tier
```

Out[84]:

	Tier	Count
0	3.0	64.08
1	1.0	35.88
2	2.0	0.04

In [85]:

```
datatmp['diff_desig'] = datatmp['ctc'] - datatmp['mean ctc']
datatmp['diff_class'] = datatmp['ctc'] - datatmp['mean2 ctc']
datatmp['diff_tier'] = datatmp['ctc'] - datatmp['mean3 ctc']
```

In [86]: `datatmp.head()`

Out[86]:

	company_hash	email_hash	orgyear	ctc
0	atrgxnnt xzaxv	6de0a4417d18ab14334c3f43397fc13b30c35149d70c05...	2016.0	1100000
1	qtrxvzwt xzegwgb rxbxnta	b0aaaf1ac138b53cb6e039ba2c3d6604a250d02d5145c10...	2018.0	4499999
2	Others	4860c670bcd48fb96c02a4b0ae3608ae6fdd98176112e9...	2015.0	2000000
3	ngpgutaxv	effdede7a2e7c2af664c8a31d9346385016128d66bbc58...	2017.0	700000
4	qxen sqghu	6ff54e709262f55cb999a1c1db8436cb2055d8f79ab520...	2017.0	1400000

Answering question based on manual clustering

1. Top 10 employees (earning more than most of the employees in the company) - Tier 1

In [87]: `datatmp[datatmp['Tier'] == 1].sort_values('diff_tier', ascending=False).head(10)[`

Out[87]:

	email_hash	ctc	mean3 ctc
68747	85e685ccaf737be77245c7bd8d06f7007e37ae8fe9a112...	12350000	6.457040e+05
103579	b69935e9b3fb05bfca232556188524cf7e0106eebbc2c2...	12500000	9.358903e+05
114980	b69935e9b3fb05bfca232556188524cf7e0106eebbc2c2...	12500000	9.358903e+05
78717	9b543e7020439a87f759ac633f4c6f2bec139bb217934f...	12500000	1.108693e+06
62313	ab48cea6068c8b0b7fbf6d152c82cc041a3f76285bf7a9...	12500000	1.254363e+06
145513	f74fc71a9d1c32af699c8e7a5ed2ea6c6fc47413f14194...	12000000	7.547604e+05
112344	f213d7959b9f75d66d69bfa430076448295d31de469378...	12300000	1.157061e+06
2628	27e22a8ce7a77250ad179c6243cf40dc4857c3026f20d8...	12000000	9.328478e+05
62255	25dda68e55433fb5c5c231e7212ad79973d1449f9f0a909...	12000000	9.527682e+05
130107	872e57a0737d0a7b5f0298586686c21c06c31fb7d5f364...	12100000	1.108693e+06

2. Top 10 employees of data science in Amazon / TCS etc earning more than their peers - Class 1

In [88]: `datatmp[(datatmp['Tier'] == 1)&(datatmp['Class'] == 1)&(datatmp['job_position'].
'ios engineer', 'data analyst', 'qa engineer',
'engineering leadership', 'data scientist', 'sdet',
'support engineer', 'security leadership', 'devops engineer']))].sort_val`

Out[88]:

			email_hash	ctc	mean2 ctc
145513	f74fc71a9d1c32af699c8e7a5ed2ea6c6fc47413f14194...		12000000	8.500900e+05	
112745	5a004387e99e281d63b1ca7f65b56fc1cdc1d1da383b58...		12000000	1.005029e+06	
84732	2bb51813fe05ab4e8ea6ba79a4c8ef631ed62f87b709ea...		12000000	1.005029e+06	
148111	f7ed6c500aac4ed1d10f8e06c85557834c966f0ac2da87...		12000000	1.005029e+06	
197421	05f1a6cd3688f5f5705e3d0752f65968a8aeb14af2261c...		12000000	1.005029e+06	
107958	d651412f3455d6fb663a5bdb84a97a75cf062ca42e4bd5...		12000000	1.005029e+06	
119289	98d1748ea1542882feb3d9075ea8a02042aaed2bfd9ecb...		12000000	1.007498e+06	
62255	25dda68e55433fbc5c231e7212ad79973d1449f9f0a909...		12000000	1.093962e+06	
38	69ef6838be8ee5b628375b4cc160ba54c1f7cab8c3b130...		11800000	9.192897e+05	
15310	02ec75722194e07fc31a7929843fc3c5fb3dd3a1e1d8b...		12500000	1.943600e+06	

3. Bottom 10 employees of data science in Amazon / TCS etc earning less than their peers - Class 3

In [89]:

```
datatmp[(datatmp['Tier'] == 1)&(datatmp['Class'] == 3)&(datatmp['job_position'].  
'ios engineer', 'data analyst', 'qa engineer',  
'engineering leadership', 'data scientist', 'sdet',  
'support engineer', 'security leadership', 'devops engineer'])).sort_val
```

Out[89]:

			email_hash	ctc	mean2 ctc
77506	74dea008e581b2b8449afdd2d694021cb3d706c667a77d...		2500000	6.250000e+06	
21889	831928c136906a9ab204fc5d793eca10d876ab2e81367a...		2500000	5.573077e+06	
49652	afb0edb04b358288621e0eb3b9544173740856ff4feb60...		2300000	5.350000e+06	
161094	e5925e18906ae284af8013ac1344b7fc185df5ace673a1...		1800000	4.700000e+06	
167535	03b1c6a500e5ced30020aac602d4aeacbe1a6b6f9181e6...		3500000	6.300000e+06	
146356	e5925e18906ae284af8013ac1344b7fc185df5ace673a1...		1800000	4.470000e+06	
76994	f3b2641868f0a42de425acdc74e2cff0d03e1a70f11af4...		3000000	5.573077e+06	
63042	1f0c6873153c336f38d3ff77b7de496a38e331bd8b1ff3...		3600000	6.156667e+06	
123281	cbbf3f54142b858711a2719a8243304ae1d5c390e449ee...		1750000	4.262500e+06	
188858	9a7894e045c744d8c5cb0b1d7281d26822183f38c62802...		3300000	5.733333e+06	

4. Bottom 10 employees (earning less than most of the employees in the company)- Tier 3

In [90]:

```
datatmp[datatmp['Tier'] == 3].sort_values('diff_tier', ascending=True).head(10)[[
```

Out[90]:

			email_hash	ctc	mean3 ctc
58997	d4450a3ecc4f69afea7a122baaf858fa08b29196c97d79...	75000	6.337000e+06		
175984	faa547efa3f271917c9429d39ee7ca1c37b48568b9ac9c...	110000	6.337000e+06		
149797	aa973bb12d00382c7e6b37833249c623cf5bd5954d6a4f...	79000	6.060375e+06		
162874	a45faf4d214fb83e7c9c64971deab8e7037433a140cbfd...	100000	6.060375e+06		
192996	6bc12bffaad2c9c9d7b098f2c2b48e7697280c714996cb...	120000	6.060375e+06		
156343	fc5ddd9f8d1bdb40fec162ffc268226a82d07527fcdb08...	120000	6.060375e+06		
175322	56bf041cf16bbf23034e5111856624e9bea0e4f06b86e8...	210000	6.060375e+06		
127722	56bf041cf16bbf23034e5111856624e9bea0e4f06b86e8...	210000	6.060375e+06		
195322	56bf041cf16bbf23034e5111856624e9bea0e4f06b86e8...	210000	6.060375e+06		
37595	e619ce4bc238af5191a8a37e484d8ca2b115c05407f2c5...	400000	4.742857e+06		

5. Top 10 employees in Amazon- X department - having 5/6/7 years of experience earning more than their peers - Tier X

In [91]: `datatmp[(datatmp['years_of_experience'].isin([5,6,7]))&(datatmp['company_hash'].isin(['...']))]`

Out[91]:

			email_hash	ctc	mean ctc
109834	2fab5e919a339803876fb532a618ab93c7b83c49746dd7...	3160000	1.760000e+06		
85760	2fab5e919a339803876fb532a618ab93c7b83c49746dd7...	1750000	1.158571e+06		
100859	5b09bddfe861cf5609982adc0a9ed3946fc08151b06cd6...	1950000	1.495000e+06		
175538	c9e14b4d46b1a76974a2e06bc546886cff85bd441f21b8...	1600000	1.158571e+06		
191007	4fb281fb2e423aa61d4cc18a09eb9fa7a68f948f6e880b...	1900000	1.555000e+06		
153368	8b3710b3a42e17677dcfbc55e48deac378af9810efa86d...	1750000	1.446667e+06		
26032	c9e14b4d46b1a76974a2e06bc546886cff85bd441f21b8...	2000000	1.760000e+06		
163141	6a3ed398f74a8186b52f98a4f6b0894beb2f2032f9af2a...	1540000	1.446667e+06		
177723	400aea75dc1316022b8c4436c60a0646fbea2962e26a5a...	1210000	1.158571e+06		
21284	30a88256b5586ba59b25e6fe78fada76950fd65ca9f250...	1200000	1.158571e+06		

6. Top 10 companies (based on their CTC)

In [92]: `datatmp.groupby('company_hash').mean()['ctc'].reset_index().sort_values('ctc', ascending=False).head(10)`

Out[92]:

	company_hash	ctc
362	bxwqgonqvntsj	6.337000e+06
3096	wvqttb	6.060375e+06
2814	vxqugqno vhnygqxnj ge xzavx	4.742857e+06
1551	orxwt	4.548000e+06
592	evzvnxwo xzw	4.383000e+06
163	bgngqi	4.133333e+06
2711	vruyvsqtu otwhqxnxt	4.009091e+06
724	gqvwr wrgha xzeqvonqhwnhqt	3.971667e+06
1207	nqvexshqv	3.960000e+06
1343	nxat	3.892000e+06

7. Top 2 positions in every company (based on their CTC)

In [93]:

```
tmp = datatmp[datatmp['job_position'] != 'na']
tmp = tmp.groupby(['company_hash', 'job_position']).mean().sort_values(['company_hash', 'job_position'])
tmp = tmp.groupby('company_hash').head(2)[['company_hash', 'job_position']]
tmp
```

Out[93]:

	company_hash	job_position
0	Others	research assistant
1	Others	researcher
269	a ntwyzgrgsxto	fullstack engineer
270	a ntwyzgrgsxto	frontend engineer
275	aaqxctz avnv owxtzwto vzvrjnxwo ucn rna	engineering intern
...
25563	zxyrtzn ntwyzgrgsxto	android engineer
25569	zxzlwwvqn	other
25570	zxzlwwvqn	area operations manager
25578	zxztrtvuo	other
25579	zxztrtvuo	engineering intern

7414 rows × 2 columns

Preparing data for training model(Imputation/Scaling)

Transforming ctc feature using log function

```
In [94]: datatmp['ctc_log'] = np.log2(datatmp['ctc'])
```

Columns like ['job_position', 'email_hash', 'company_hash'] are text.

We can't use them during imputation, so we'll remove these columns

```
In [95]: drop_cols = ['job_position', 'email_hash', 'company_hash']
for i in drop_cols:
    try:
        datatmp.drop([i], axis=1, inplace=True)
    except:
        print('no')
```

```
In [96]: datatmp.columns
```

```
Out[96]: Index(['orgyear', 'ctc', 'ctc_updated_year', 'years_of_experience', 'mean ctc',
       'Designation', 'mean2 ctc', 'Class', 'mean3 ctc', 'Tier', 'diff_desig',
       'diff_class', 'diff_tier', 'ctc_log'],
      dtype='object')
```

```
In [97]: datatmp.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 201413 entries, 0 to 201412
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   orgyear          201413 non-null   float64 
 1   ctc              201413 non-null   int64  
 2   ctc_updated_year 201413 non-null   float64 
 3   years_of_experience 201413 non-null   float64 
 4   mean ctc         201413 non-null   float64 
 5   Designation       201413 non-null   float64 
 6   mean2 ctc         201413 non-null   float64 
 7   Class             201413 non-null   float64 
 8   mean3 ctc         201413 non-null   float64 
 9   Tier              201413 non-null   float64 
 10  diff_desig        201413 non-null   float64 
 11  diff_class        201413 non-null   float64 
 12  diff_tier         201413 non-null   float64 
 13  ctc_log           201413 non-null   float64 
dtypes: float64(13), int64(1)
memory usage: 23.0 MB
```

Statistical Summary

```
In [98]: datatmp.describe()
```

Out[98]:

	orgyear	ctc	ctc_updated_year	years_of_experience	mean ctc
count	201413.000000	2.014130e+05	201413.000000	201413.000000	2.014130e+05
mean	2015.103722	1.313212e+06	2019.687314	8.896278	1.313212e+06
std	4.257499	1.234111e+06	1.287119	4.257499	9.428384e+05
min	1990.000000	3.800000e+04	2015.000000	2.000000	3.900000e+04
25%	2013.000000	5.500000e+05	2019.000000	6.000000	7.195652e+05
50%	2016.000000	9.500000e+05	2020.000000	8.000000	1.034220e+06
75%	2018.000000	1.650000e+06	2021.000000	11.000000	1.600000e+06
max	2022.000000	1.250000e+07	2022.000000	34.000000	1.250000e+07

In [99]: `datatmp.isna().sum()`

Out[99]:

orgyear	0
ctc	0
ctc_updated_year	0
years_of_experience	0
mean ctc	0
Designation	0
mean2 ctc	0
Class	0
mean3 ctc	0
Tier	0
diff_desig	0
diff_class	0
diff_tier	0
ctc_log	0
dtype: int64	

Training Model

In [100...]:

```
from sklearn.impute import KNNImputer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.cluster import MiniBatchKMeans, KMeans
from sklearn.metrics import silhouette_score
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
```

Standardization

In [101...]:

```
scaler = StandardScaler()
scaler.fit(datatmp)
X = pd.DataFrame(scaler.transform(datatmp), columns=datatmp.columns, index=datatmp.index)
```

In [102...]:

X

Out[102]:

	orgyear	ctc	ctc_updated_year	years_of_experience	mean ctc	Designa
0	0.210518	-0.172766		0.242935	-0.210518	-0.226139
1	0.680278	-0.699463		-0.533996	-0.680278	-0.571601
2	-0.024362	0.556506		0.242935	0.024362	-0.131941
3	0.445398	-0.496887		-0.533996	-0.445398	-0.164016
4	0.445398	0.070325		-0.533996	-0.445398	0.092050
...
201408	-1.668524	-0.885832		-0.533996	1.668524	-1.159493
201409	0.445398	-0.658947		0.242935	-0.445398	-0.173107
201410	1.384919	-0.496887		1.019866	-1.384919	-0.685745
201411	0.915159	3.068442		-0.533996	-0.915159	-0.116674
201412	-0.259243	-0.059324		-2.864788	0.259243	0.403168

201413 rows × 14 columns

Hierarchical Clustering

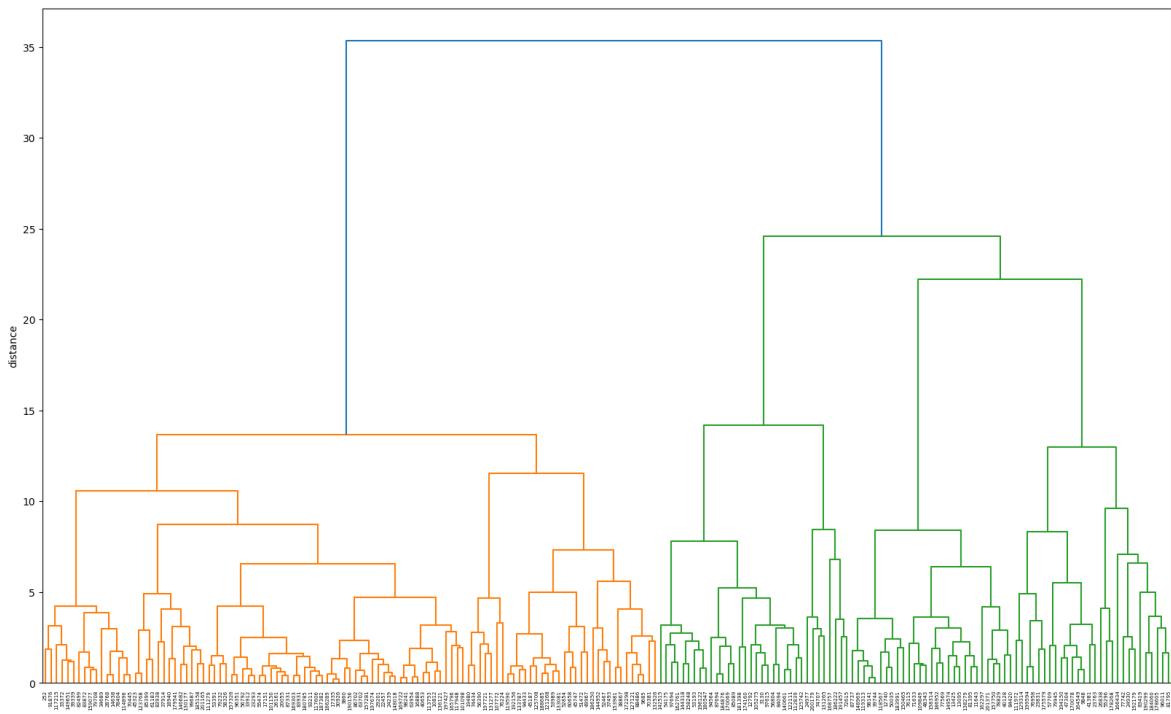
Trying to get a high level idea about how many clusters we can from, by taking sample of 200 learners multiple times and forming hierarchy and visualising in dendrogram.

In [103...]

```
import scipy.cluster.hierarchy as sch

sample = X.sample(200)
Z = sch.linkage(sample, method='ward')

fig, ax1 = plt.subplots(figsize=(20, 12))
sch.dendrogram(Z, labels=sample.index, ax=ax1)
plt.xticks(rotation=90)
ax1.set_ylabel('distance')
plt.show()
```

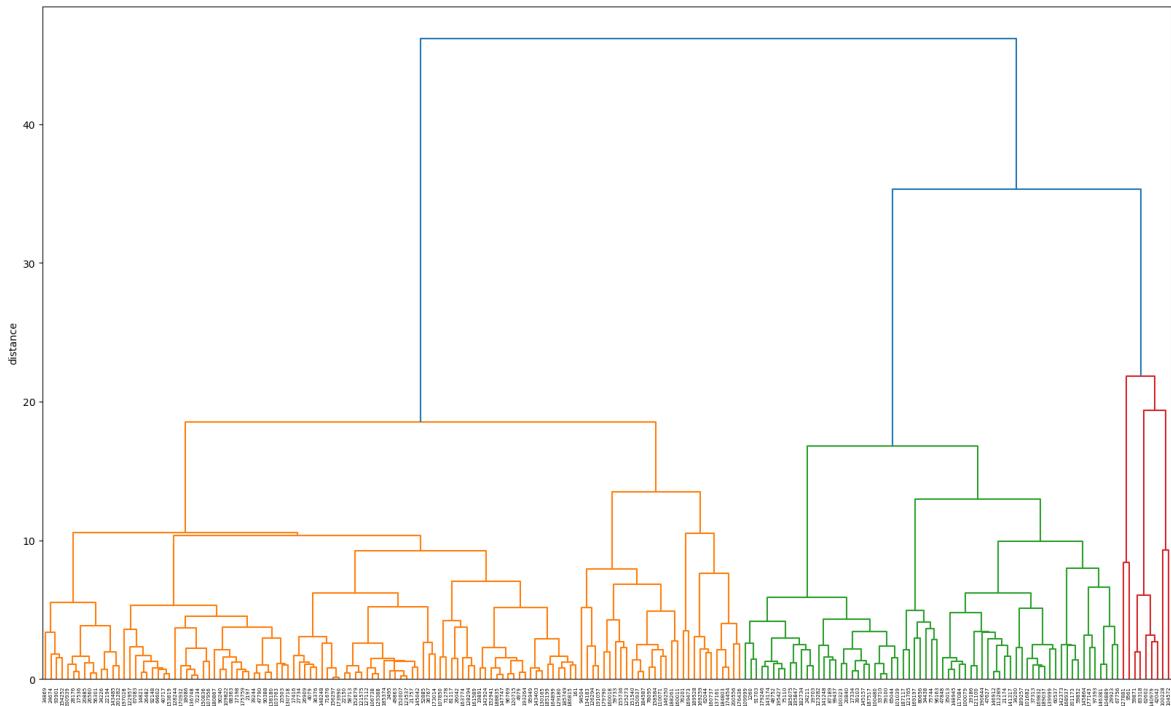


In [104]:

```
import scipy.cluster.hierarchy as sch

sample = X.sample(200)
Z = sch.linkage(sample, method='ward')

fig, ax2 = plt.subplots(figsize=(20, 12))
sch.dendrogram(Z, labels=sample.index, ax=ax2)
plt.xticks(rotation=90)
ax2.set_ylabel('distance')
plt.show()
```



In [105]:

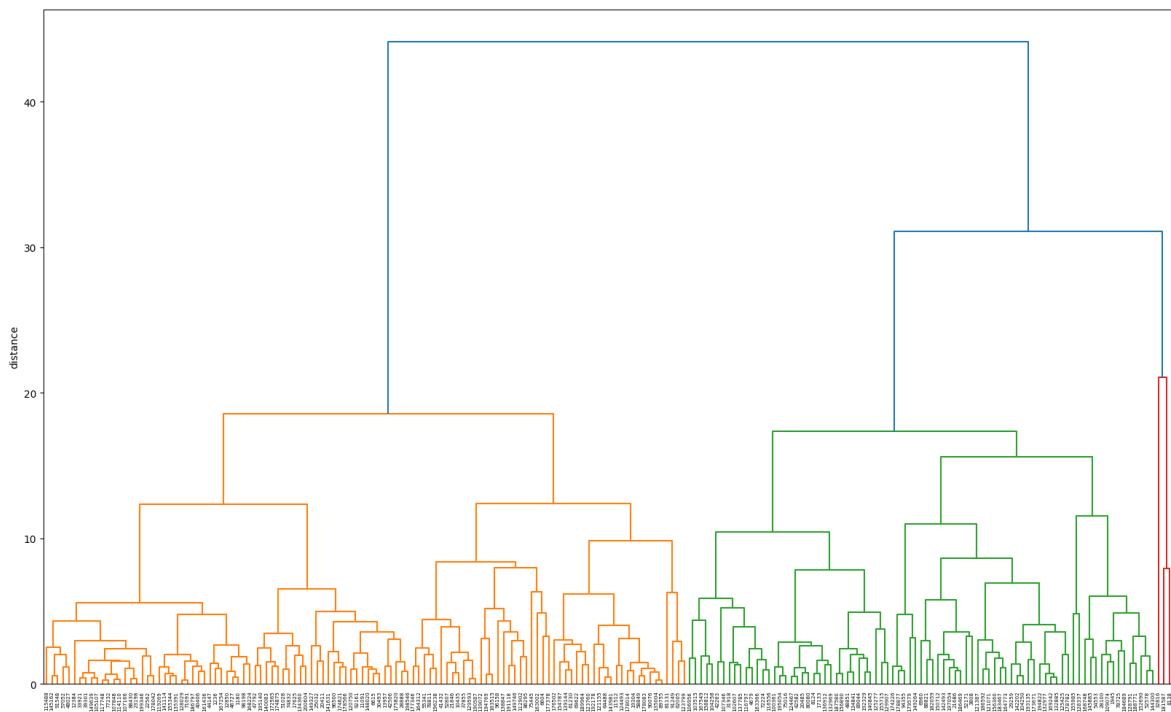
```
import scipy.cluster.hierarchy as sch

sample = X.sample(200)
Z = sch.linkage(sample, method='ward')
```

```

fig, ax3 = plt.subplots(figsize=(20, 12))
sch.dendrogram(Z, labels=sample.index, ax=ax3)
plt.xticks(rotation=90)
ax3.set_ylabel('distance')
plt.show()

```



- Based on dendrogram, we can observe there are 3 clusters in the data based on similarity

Further checking appropriate number of clusters using Elbow Method using k-Means clustering :

Kmeans clustering

```

In [106... pipe_knn = Pipeline([('scaler', StandardScaler()), ('knn_imputer', KNNImputer(n
In [107... pipe_knn_5 = Pipeline([('scaler', StandardScaler()), ('knn_imputer', KNNImputer(
In [108... pipe = Pipeline([('scaler', StandardScaler()), ('simple_imputer', SimpleImputer(
In [109... pipe_knn_pca = Pipeline([('scaler', StandardScaler()), ('knn_imputer', KNNImput
In [110... pipe_unscaled = Pipeline([('knn_imputer', KNNImputer(n_neighbors=5, weights="un

```

Finding optimal num of clusters using Elbow method

```

In [111... for name,pipeline in [('KNN Imputation',pipe_knn),('KNN Imputation with (defaul
                           ('KNN Imputation + PCA', pipe_knn_pca),('KNN Imputation U
                           X = pipeline.fit_transform(datatmp)
                           X = pd.DataFrame(X)

```

```

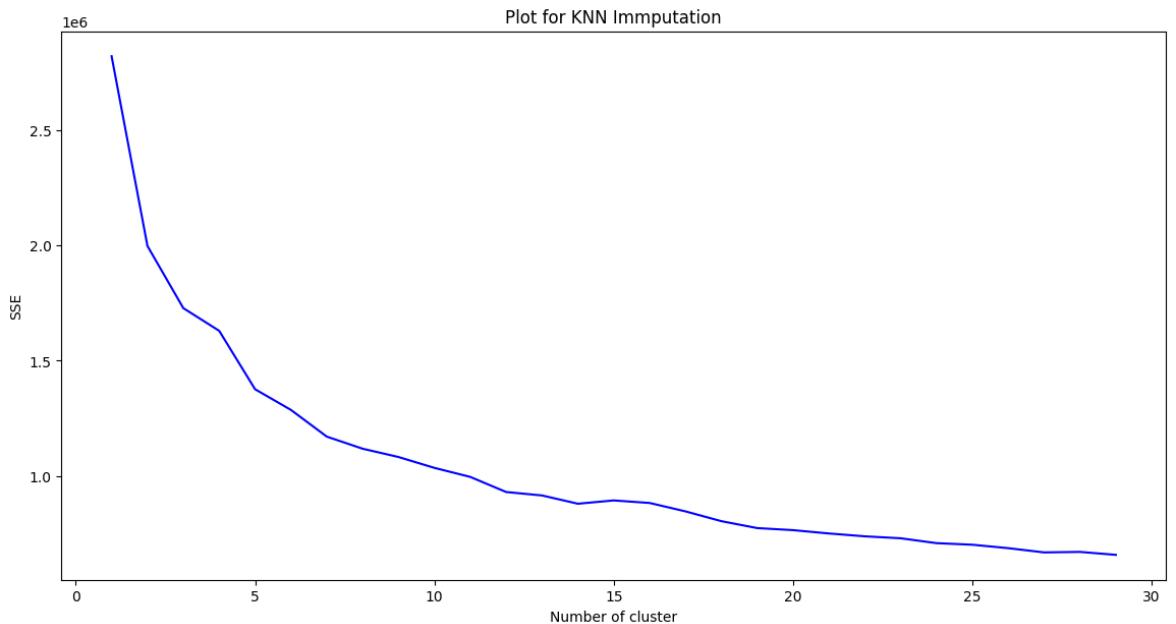
if "PCA" not in name :
    X.columns= datatmp.columns

sse = {}
print("Running for ",name)
for k in range(1, 30):
    kmeans = MiniBatchKMeans(init="k-means++",n_clusters=k,
                             random_state=0).fit(X)
    label = kmeans.labels_
    datatmp["clusters"] = label
    sse[k] = kmeans.inertia_

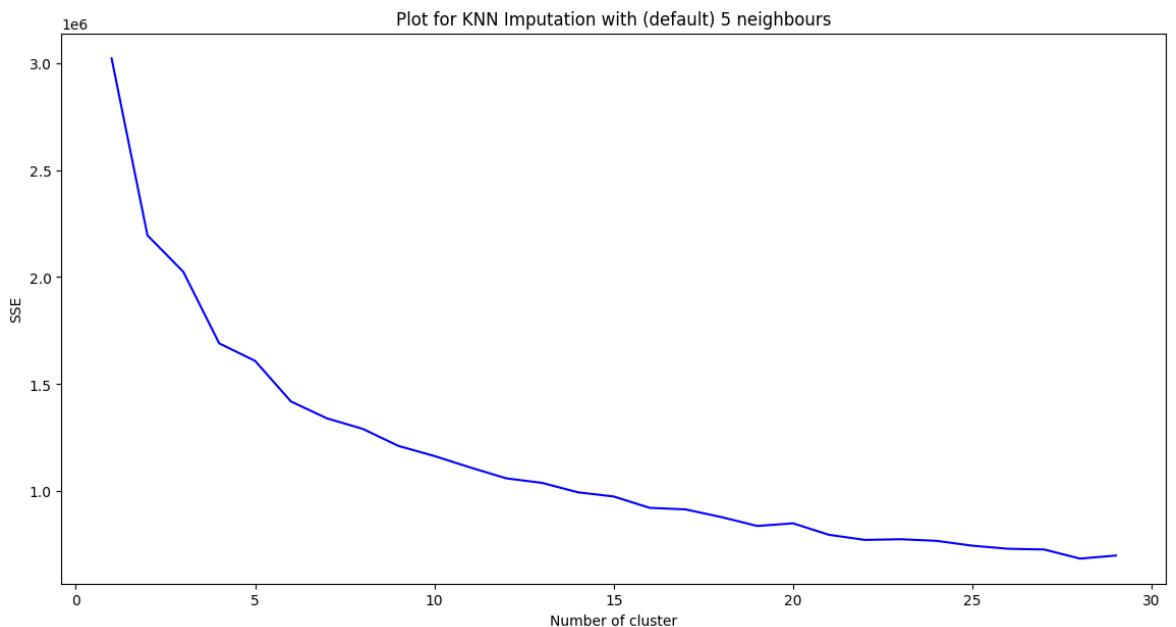

plt.figure(figsize=(14,7))
plt.plot(list(sse.keys()), list(sse.values()),'b-',label='Sum of squared err
plt.xlabel("Number of cluster")
plt.ylabel("SSE")
plt.title("Plot for "+name)
plt.show()

```

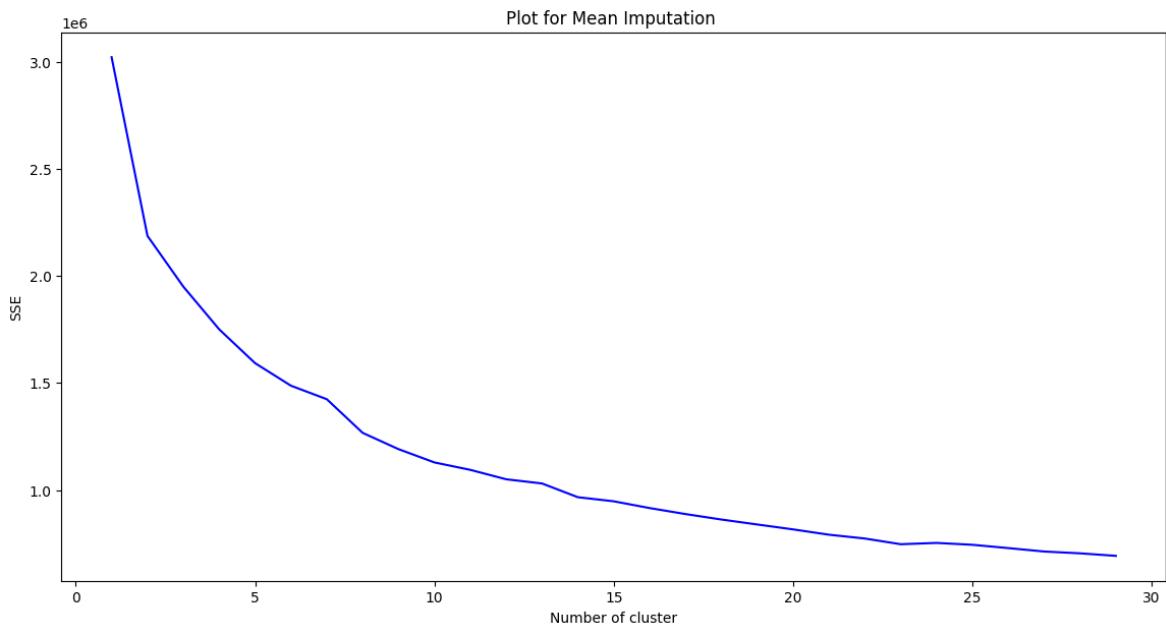
Running for KNN Imputation



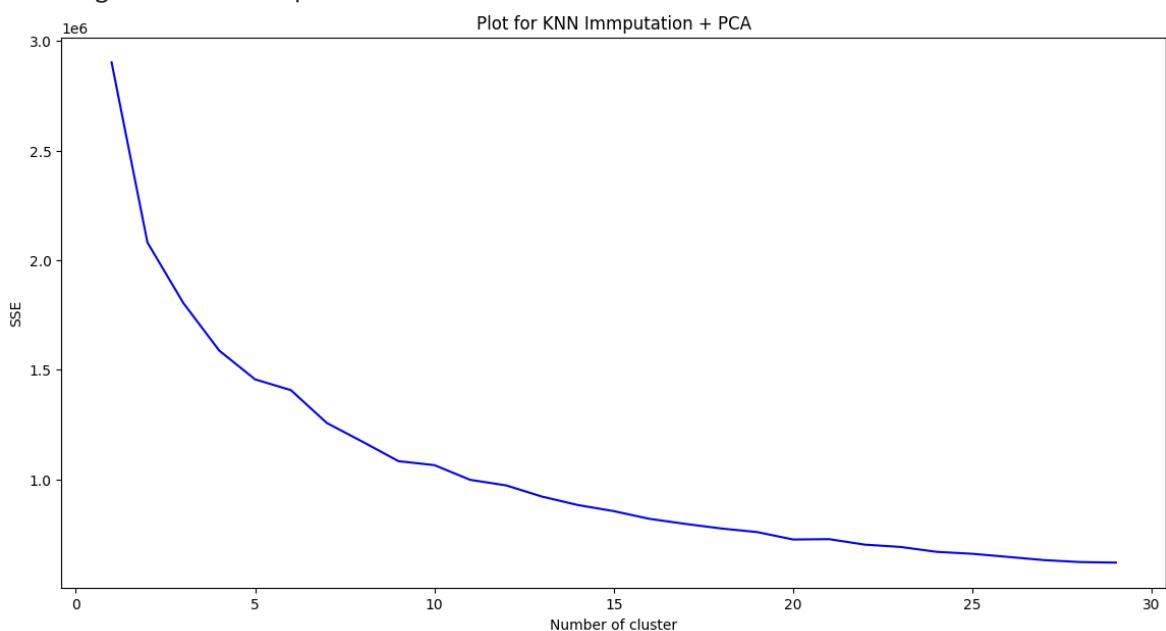
Running for KNN Imputation with (default) 5 neighbours



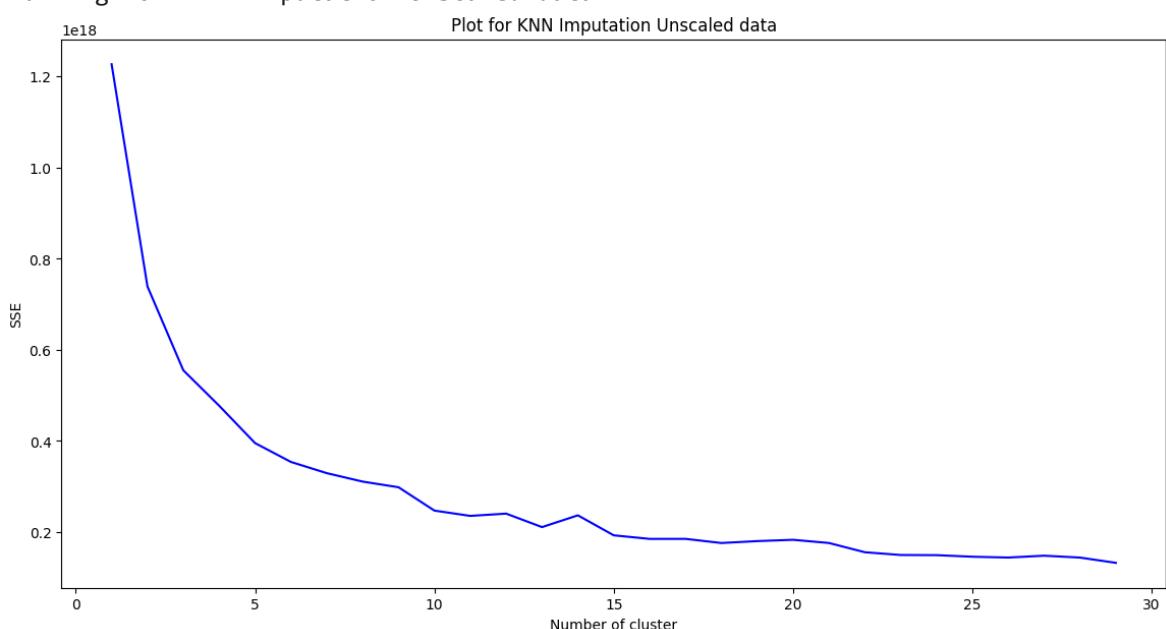
Running for Mean Imputation



Running for KNN Imputation + PCA



Running for KNN Imputation Unscaled data



Model	N_clusters
KNN Imputation	16
KNN Imputation with (default) 5 neighbours	20
Mean Imputation	25
KNN Imputation + PCA	21
KNN Imputation Unscaled data	5

- Number of clusters is around 18-20 for scaled data, while around 5 for unscaled data

```
In [112...]: for i in range(1,10):
    k = 4

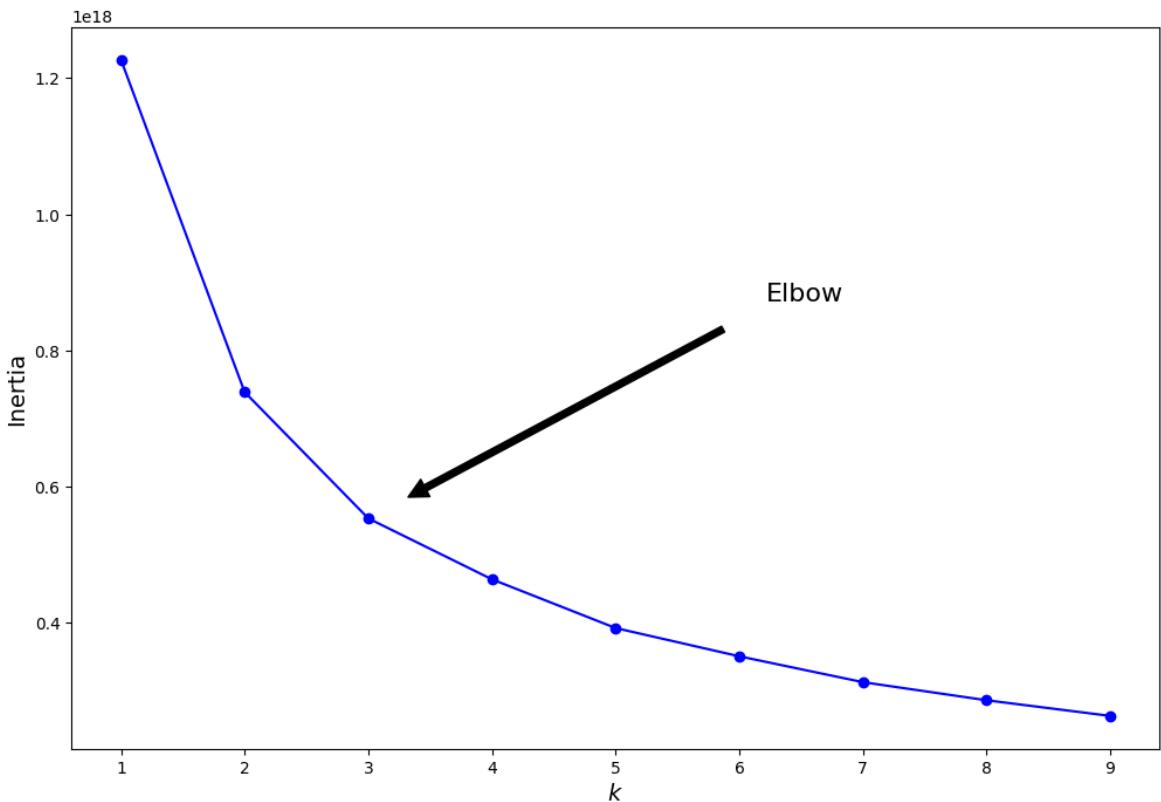
    kM = KMeans(n_clusters=k,
                 random_state=654)
    y_pred = kM.fit_predict(X)
```

```
In [113...]: kmeans_per_k = [KMeans(n_clusters=k, random_state=42).fit(X)
                     for k in range(1, 10)]

inertias = [model.inertia_ for model in kmeans_per_k]
inertias
```

Out[113]: [1.2270254261284096e+18,
 7.387574640506063e+17,
 5.532265700854374e+17,
 4.637107637866577e+17,
 3.9211306753527846e+17,
 3.5060402966441926e+17,
 3.126052869164911e+17,
 2.8600457163032083e+17,
 2.6300352325099155e+17]

```
In [114...]: plt.figure(figsize=(12, 8))
plt.plot(range(1, 10), inertias, "bo-")
plt.xlabel("$k$", fontsize=14)
plt.ylabel("Inertia", fontsize=14)
plt.annotate('Elbow',
             xy=(3, inertias[2]),
             xytext=(0.55, 0.55),
             textcoords='figure fraction',
             fontsize=16,
             arrowprops=dict(facecolor='black', shrink=0.1)
            )
plt.show()
```

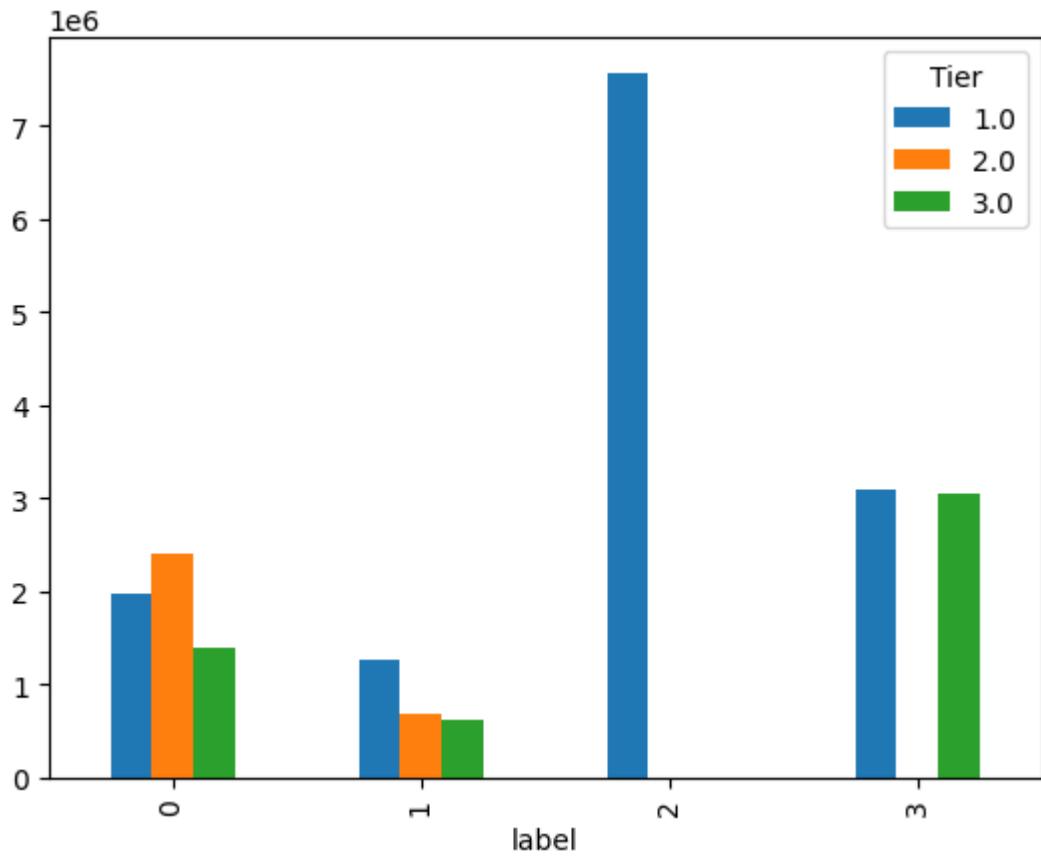


```
In [115]: clusters = pd.DataFrame(X, columns=X.columns)
clusters['label'] = kM.labels_
clusters.sample(5)
```

	orgyear	ctc	ctc_updated_year	years_of_experience	mean ctc	Design
171635	2012.0	742000.0		2016.0	12.0	1.794000e+06
5086	2014.0	1000000.0		2021.0	10.0	9.300000e+05
93005	2022.0	300000.0		2022.0	2.0	5.250000e+05
89476	2020.0	400000.0		2021.0	4.0	4.438887e+05
69628	2012.0	2111000.0		2017.0	12.0	2.076200e+06

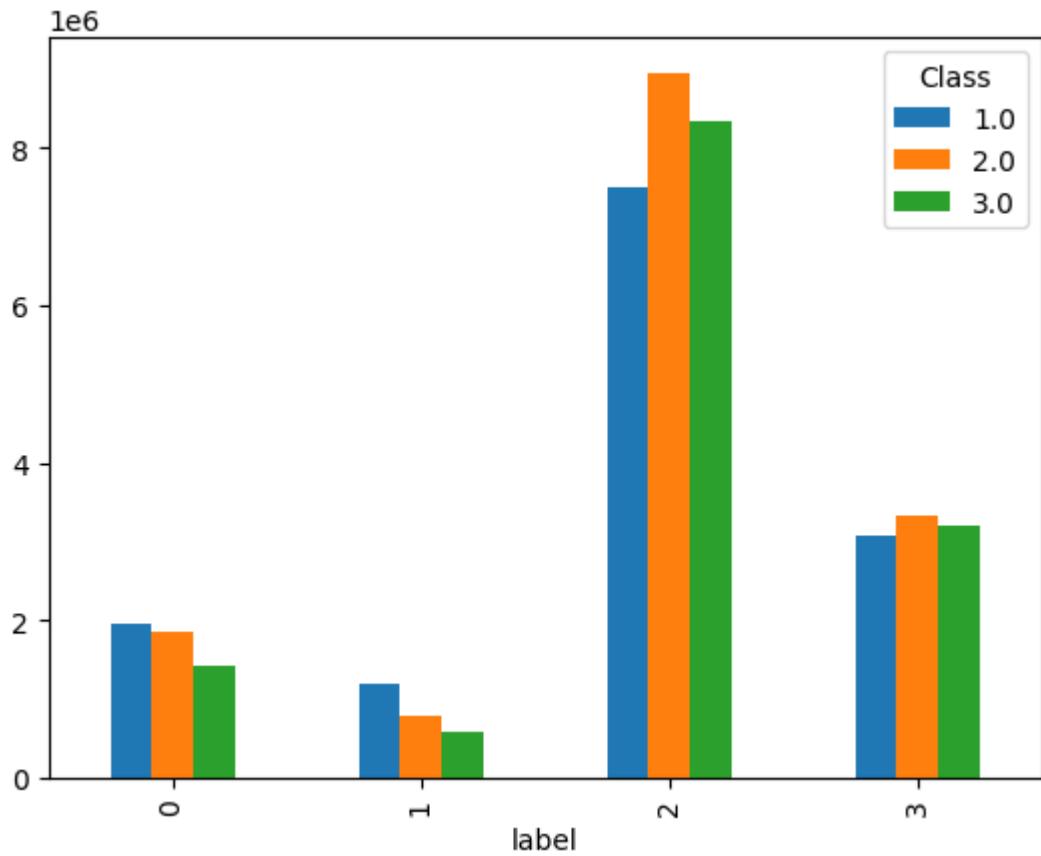
EDA after Clustering

```
In [116]: pd.crosstab(index = clusters["label"],
                    columns = clusters["Tier"], values=clusters["ctc"], aggfunc= np.mean
                  ).plot(kind = "bar")
plt.show()
```

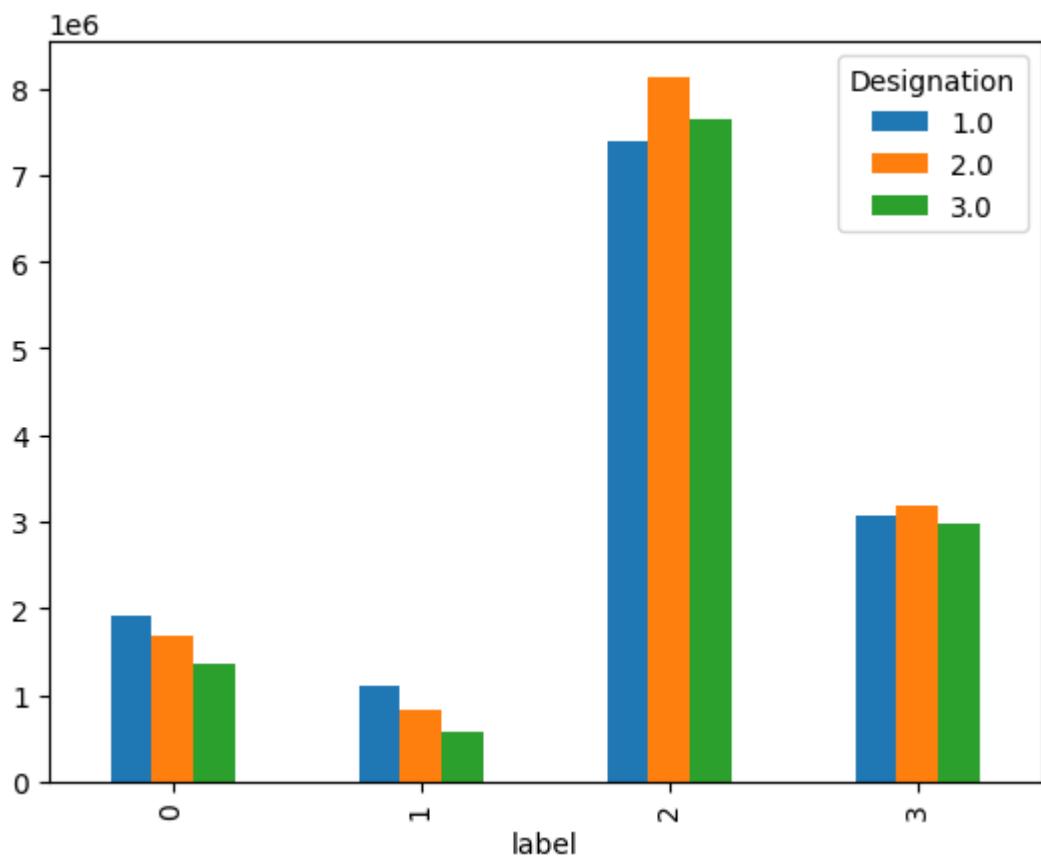


- Based on k-Means Clustering algorithm output , as well as manual clustering , learners from tier1 company receiving very high CTC.

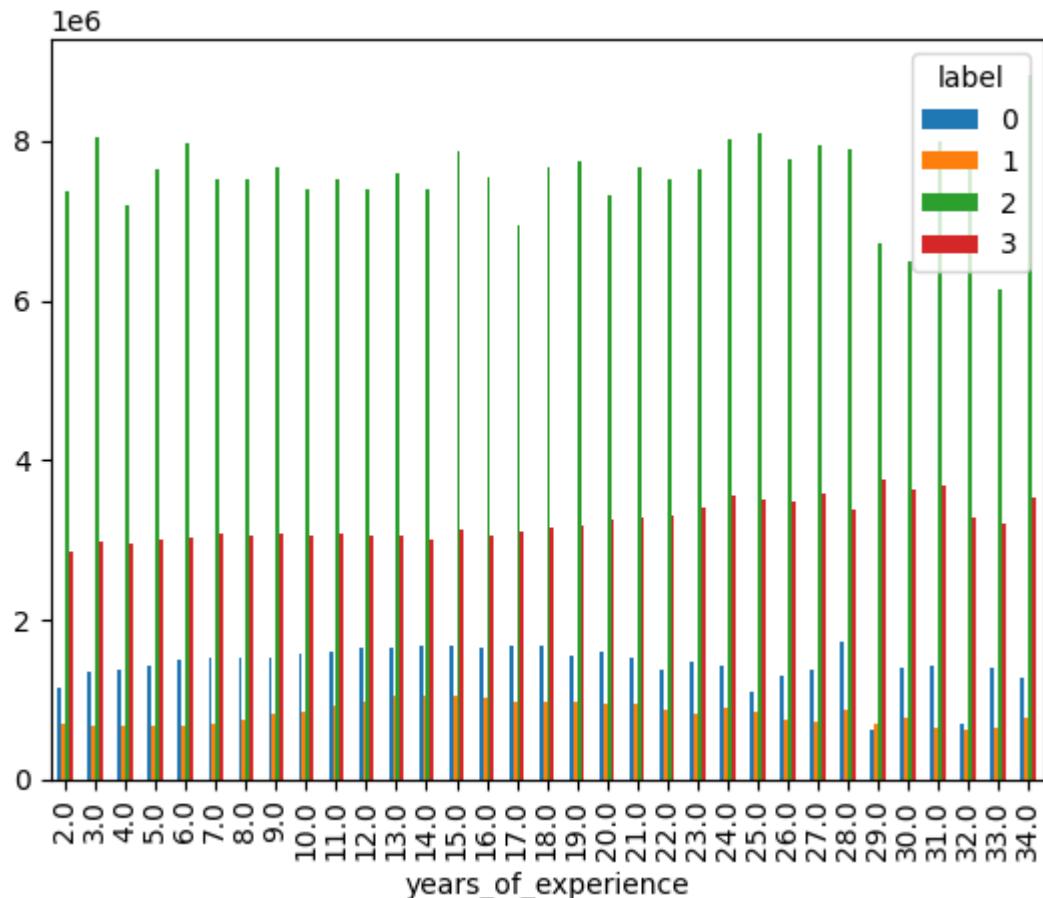
```
In [117]: pd.crosstab(index = clusters["label"],
                  columns = clusters["Class"], values=clusters["ctc"], aggfunc= np.mean
                  ).plot(kind = "bar")
plt.show()
```



```
In [118]: pd.crosstab(index = clusters["label"],
                   columns = clusters["Designation"],
                   values=clusters["ctc"],aggfunc= np.mean
                  ).plot(kind = "bar")
plt.show()
```



```
In [119...]: pd.crosstab(columns = clusters["label"],
    index = clusters["years_of_experience"],
    values=clusters["ctc"],aggfunc= np.mean
    ).plot(kind = "bar")
plt.show()
```

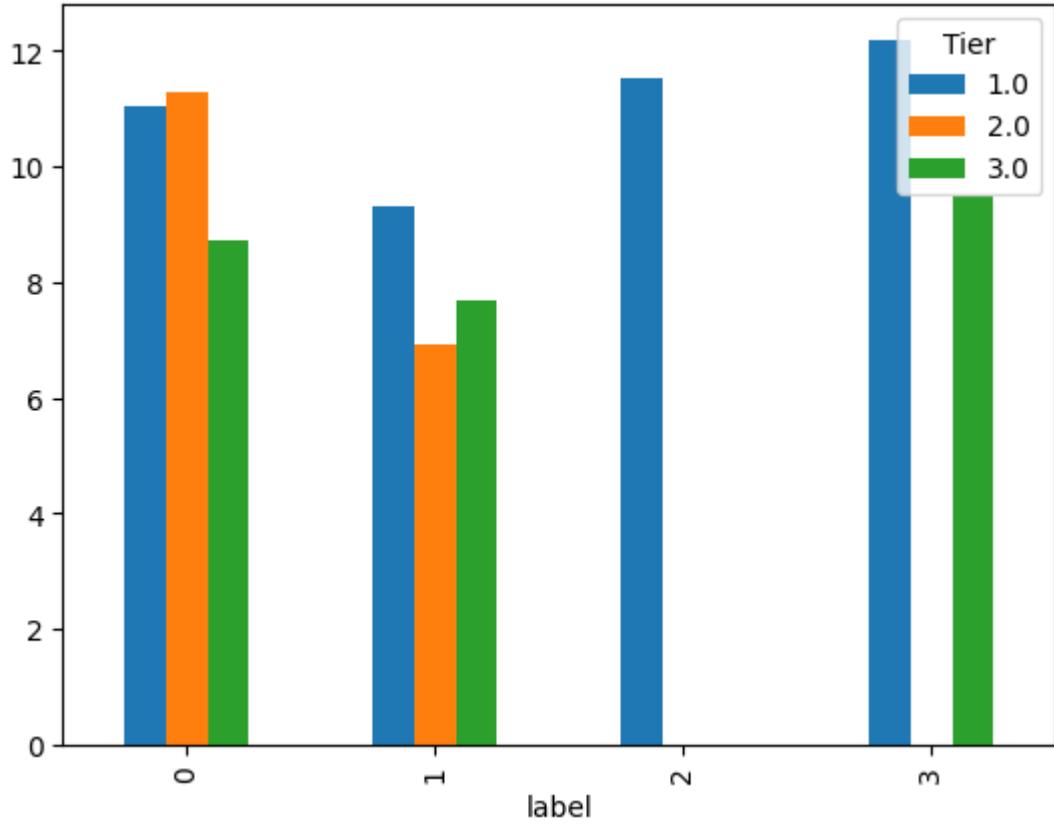


- Cluster label 2, are those learners who are very very experienced.

```
In [120...]: clusters.label.value_counts(normalize=True)*100
```

```
Out[120]: 1    68.080015
0    16.797327
3    13.748864
2    1.373794
Name: label, dtype: float64
```

```
In [121...]: # years_of_experience per each cluster group of Learners
pd.crosstab(index = clusters["label"],
    columns = clusters["Tier"],
    values=clusters["years_of_experience"],
    aggfunc=np.mean
    ).plot(kind = "bar")
plt.show()
```



Statistical Summary based on Each Cluster

```
In [122]: clusters.groupby("label").describe()[[ "ctc", "Class", "Tier", "years_of_experience"]]
```

Out[122]:

	label	0	1	2	3			
ctc	count	3.383200e+04	1.371220e+05	2.767000e+03	2.769200e+04			
	mean	1.539786e+06	7.713451e+05	7.564227e+06	3.094949e+06			
	std	6.210890e+05	4.148623e+05	1.942410e+06	8.840100e+05			
	min	3.955000e+04	3.800000e+04	4.400000e+06	1.700000e+06			
	25%	1.200000e+06	4.500000e+05	6.000000e+06	2.400000e+06			
	50%	1.600000e+06	7.000000e+05	7.000000e+06	2.960000e+06			
	75%	2.000000e+06	1.000000e+06	9.000000e+06	3.600000e+06			
	max	4.000000e+06	2.200000e+06	1.250000e+07	7.500000e+06			
	Class	count	3.383200e+04	1.371220e+05	2.767000e+03	2.769200e+04		
		mean	2.578092e+00	2.376577e+00	1.043730e+00	1.139715e+00		
		std	7.785757e-01	8.921285e-01	2.114823e-01	4.376737e-01		
		min	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00		
		25%	3.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00		
		50%	3.000000e+00	3.000000e+00	1.000000e+00	1.000000e+00		
		75%	3.000000e+00	3.000000e+00	1.000000e+00	1.000000e+00		
		max	3.000000e+00	3.000000e+00	3.000000e+00	3.000000e+00		
		Tier	count	3.383200e+04	1.371220e+05	2.767000e+03	2.769200e+04	
			mean	2.529233e+00	2.504609e+00	1.000000e+00	1.005778e+00	
			std	8.481931e-01	8.630596e-01	0.000000e+00	1.073439e-01	
			min	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	
			25%	3.000000e+00	3.000000e+00	1.000000e+00	1.000000e+00	
			50%	3.000000e+00	3.000000e+00	1.000000e+00	1.000000e+00	
			75%	3.000000e+00	3.000000e+00	1.000000e+00	1.000000e+00	
			max	3.000000e+00	3.000000e+00	1.000000e+00	3.000000e+00	
			years_of_experience	count	3.383200e+04	1.371220e+05	2.767000e+03	2.769200e+04
				mean	9.260877e+00	8.089395e+00	1.153993e+01	1.218211e+01
				std	3.832685e+00	3.600889e+00	5.868754e+00	5.600577e+00
				min	2.000000e+00	2.000000e+00	2.000000e+00	2.000000e+00
				25%	6.000000e+00	6.000000e+00	7.000000e+00	8.000000e+00
				50%	9.000000e+00	7.000000e+00	1.000000e+01	1.100000e+01
				75%	1.100000e+01	1.000000e+01	1.500000e+01	1.500000e+01
				max	3.400000e+01	3.400000e+01	3.400000e+01	3.400000e+01



Insights

- Top Paying job titles include full stack engineer', 'iOS engineer', 'data analyst', 'qi engineer', 'engineering leadership', 'data scientist', 'sdet', 'support engineer', 'security leadership', 'devops engineer'.
- Among top paying companies, salary for these is getting lesser in recent years, Goldmaan Sachs, Tata Consultancy Services, Samsung Electronics, VMware, Dell, Dbs Bank, Hsbc software development India and GE
- Among Top paying companies mean salary for these company is increasing every year, Amazon, Microsoft and Huawei Technologies
- Avg. CTC seems to be decreasing with year.
- 1017 unique job positions are there in the dataset.
- 857 unique job positions are there in the dataset after pre-processing strings.
- Number of clusters is around 18-20 for scaled data, while around 5 for unscaled data.
- Majority number of Tier is 3.
- As compare to other Cluster label 1 is the high value.
- Based on k-Means Clustering algorithm output , as well as manual clustering , learners from tier1 company receiving very high CTC.
- Cluster label 2, are those learners who are very very experienced.
- Based on dendrogram , we can observe there are 3 clusters in the data based on similarity.

Recommendations

- Freshers who want to work on technical side should look for roles related to Backend Engineer, SDET, QA engineer, Data Scientist, Android Engineer, Full stack engineer to get good salaries as experience increases.
- Freshers who want best CTC should aim for companies like 'Cisco', 'Intel Technology India Pvt Ltd', 'Amazon', 'Walmart Labs', 'Symantec', 'Schneider Electric India', 'Morgan Stanley', 'Ericsson RD Bangalore' and 'Samsung Electronics'.